

PA1_na21b033

September 10, 2023

#1. Self Implementation

##1.1. Installing necessary packages

```
[38]: !pip install -q torch_snippets
      !pip install -q torchinfo
      !pip install -q torchmetrics

from torchmetrics import ConfusionMatrix
from torch_snippets import*
from torchinfo import summary
import torch
from torchvision import datasets, transforms
from torch import nn
import matplotlib.pyplot as plt
import sklearn
from sklearn.metrics import classification_report
import pandas as pd
```

```
[39]: device="cuda" if torch.cuda.is_available() else "cpu"
      device #device agnostic code
```

```
[39]: 'cuda'
```

##1.2. Getting the data

```
[40]: transformer= transforms.Compose([transforms.ToTensor()]) #dataset contains imgs
      ↪in PIL format
```

```
[41]: train_ds=datasets.MNIST(root="MNIST/", train=True, download=True, transform=
      ↪transformer, target_transform=torchvision.transforms.Compose([lambda x:torch.
      ↪LongTensor([x]), lambda x:F.one_hot(x,10)]))
      #convert labels to one hot encoding
      test_ds=datasets.MNIST(root="MNIST/", train=False, download=True, transform=
      ↪transformer, target_transform=torchvision.transforms.Compose([lambda x:torch.
      ↪LongTensor([x]), lambda x:F.one_hot(x,10)]))

      train_dl= torch.utils.data.DataLoader(train_ds, batch_size=64, drop_last=True)
      test_dl= torch.utils.data.DataLoader(test_ds, batch_size=64, drop_last=True)
```

```
[42]: im, label=next(iter(train_dl))
      im.shape, label.shape #labels are not one hot encoded; img_shape-> [28*28]
```

```
[42]: (torch.Size([64, 1, 28, 28]), torch.Size([64, 1, 10]))
```

```
[43]: im=im.view([64,-1])
      im.shape, label.shape
```

```
[43]: (torch.Size([64, 784]), torch.Size([64, 1, 10]))
```

##1.3. Building the model

```
[44]: class Layer:
      def __init__(self):
          self.input = None
          self.output = None

      def forward_propagation(self, input):
          raise NotImplementedError

      def backward_propagation(self, output_error, learning_rate):
          raise NotImplementedError
```

```
[45]: class Linear_1(Layer):
      def __init__(self, in_features, out_features):
          self.m1=(6/(in_features+out_features))*0.5
          self.weights= torch.FloatTensor(in_features, out_features).uniform_(-self.
      ↪m1, self.m1).to(device) #glorot initialization
          self.bias= torch.zeros(1, out_features).to(device)

      def forward_prop(self, input_data):
          self.input=input_data
          self.output= torch.matmul(self.input, self.weights)+self.bias
          return self.output

      def back_prop(self, output_error, lr):
          weights_error= torch.matmul(self.input.T, output_error)
          self.weights= self.weights-lr*weights_error
          self.bias= self.bias-lr*output_error
          input_error= torch.matmul(output_error, self.weights.T)
          return input_error
```

```
[46]: class Activation_1(Layer):
      def __init__(self, activation, activation_g): #activation_g-> activation_
      ↪gradient
          self.activation=activation
          self.activation_g=activation_g
```

```

def forward_prop(self, input_data):
    self.input=input_data
    return self.activation(self.input)

def back_prop(self, output_error, lr):
    return torch.mul(self.activation_g(self.input), output_error)

```

[47]: *'''Activation function ->g indicates gradient || derivative of the function'''*

```

def relu(x):
    return nn.functional.relu(x)

def relu_g(x):
    x1=x.clone()
    x1[x<=0]=0
    x1[x>0]=1
    return x1

def tanh(x):
    return nn.functional.tanh(x)

def tanh_g(x):
    return 1-(tanh(x))**2

def softmax(x):
    return nn.functional.softmax(x, dim=-1)

def softmax_g(x):
    output= softmax(x)*(1-softmax(x))
    return output

''' From experimentation it has been found that relu serves as the better
↪activation this is because of the vanishing gradients issue of the tanh
↪activation function'''

```

[47]: ' From experimentation it has been found that relu serves as the better activation this is because of the vanishing gradients issue of the tanh activation function'

```

[48]: def accuracy(y_t, y_p):
    y_t, y_p= torch.argmax(y_t, dim=-1), torch.argmax(y_p, dim=-1)
    return torch.sum(torch.eq(y_t, y_p))/len(y_t)

```

[49]: *''' Loss function and its gradient'''*

```

def cross_entropy(y_t, y_p):

```

```

        output=-torch.log(torch.sum(torch.mul(y_p, y_t), dim=1))
        return torch.sum(output)/len(y_t)

def cross_entropy_g(y_t, y_p):
    return -2*(y_t-y_p)/len(y_t)

```

```

[50]: class network:
        def __init__(self):
            self.layers=[]
            self.loss=None
            self.loss_g=None

        def add(self, layer):
            self.layers.append(layer)

        def use(self, loss, loss_g):
            self.loss= loss
            self.loss_g= loss_g

        def train_step(self, data, lr):
            x_t, y_t= data
            x_t, y_t=x_t.view([64,-1]).to(device), y_t.to(device).squeeze()
            output= x_t
            for layer in self.layers:
                output= layer.forward_prop(output)
            err=self.loss(y_t, output)
            error=self.loss_g(y_t, output)
            for layer in reversed(self.layers):
                error= layer.back_prop(error, lr)
            train_acc=accuracy(y_t, output)
            return err, train_acc

        def test_step(self, data):
            x, y= data
            x, y= x.view([64, -1]).to(device), y.to(device).squeeze()
            output= x
            for layer in self.layers:
                output= layer.forward_prop(output)
            err=self.loss(y, output)
            test_acc=accuracy(y, output)
            return err, test_acc, output

```

```

[51]: net=network()
        net.add(Linear_1(28*28, 500))
        net.add(Activation_1(relu, relu_g))
        net.add(Linear_1(500, 250))
        net.add(Activation_1(relu, relu_g))

```

```

net.add(Linear_l(250, 100))
net.add(Activation_l(relu, relu_g))
net.add(Linear_l(100, 10))
net.add(Activation_l(softmax, softmax_g))

net.use(cross_entropy, cross_entropy_g)

```

```

[52]: EPOCHS=15
log=Report(EPOCHS)
for epoch in range(EPOCHS):
    n=len(train_dl)
    for idx, data in enumerate(train_dl):
        loss, t_acc= net.train_step(data, lr=0.01)
        log.record((epoch+(1+idx)/n), train_loss= loss, train_acc= t_acc, end="\r")

    n=len(test_dl)
    for idx, data in enumerate(test_dl):
        loss, t_acc, output= net.test_step(data)
        log.record((epoch+(1+idx)/n), test_loss= loss, test_acc= t_acc, end="\r")
    log.report_avgs(epoch+1)

```

```

EPOCH: 1.000 test_loss: 2.986 test_acc: 0.768 train_acc: 0.581 train_loss:
1.946 (15.42s - 215.81s remaining)
EPOCH: 2.000 test_loss: 1.193 test_acc: 0.815 train_acc: 0.797 train_loss:
2.508 (36.60s - 237.89s remaining)
EPOCH: 3.000 test_loss: 0.604 test_acc: 0.879 train_acc: 0.854 train_loss:
0.860 (48.26s - 193.03s remaining)
EPOCH: 4.000 test_loss: 0.449 test_acc: 0.898 train_acc: 0.889 train_loss:
0.540 (60.28s - 165.76s remaining)
EPOCH: 5.000 test_loss: 0.369 test_acc: 0.909 train_acc: 0.902 train_loss:
0.429 (72.14s - 144.28s remaining)
EPOCH: 6.000 test_loss: 0.325 test_acc: 0.916 train_acc: 0.911 train_loss:
0.364 (84.49s - 126.73s remaining)
EPOCH: 7.000 test_loss: 0.299 test_acc: 0.921 train_acc: 0.916 train_loss:
0.327 (96.97s - 110.82s remaining)
EPOCH: 8.000 test_loss: 0.282 test_acc: 0.925 train_acc: 0.921 train_loss:
0.305 (110.00s - 96.25s remaining)
EPOCH: 9.000 test_loss: 0.268 test_acc: 0.928 train_acc: 0.925 train_loss:
0.288 (121.95s - 81.30s remaining)
EPOCH: 10.000 test_loss: 0.257 test_acc: 0.930 train_acc: 0.929 train_loss:
0.274 (133.88s - 66.94s remaining)
EPOCH: 11.000 test_loss: 0.247 test_acc: 0.932 train_acc: 0.932 train_loss:
0.262 (145.80s - 53.02s remaining)
EPOCH: 12.000 test_loss: 0.238 test_acc: 0.935 train_acc: 0.934 train_loss:
0.251 (162.35s - 40.59s remaining)
EPOCH: 13.000 test_loss: 0.229 test_acc: 0.938 train_acc: 0.937 train_loss:
0.242 (188.94s - 29.07s remaining)

```

```

EPOCH: 14.000 test_loss: 0.222 test_acc: 0.939 train_acc: 0.939 train_loss:
0.233 (200.95s - 14.35s remaining)
EPOCH: 15.000 test_loss: 0.216 test_acc: 0.941 train_acc: 0.941 train_loss:
0.225 (213.06s - 0.00s remaining)

```

```


[53]: fig, ax= plt.subplots(ncols=2, figsize=(15,5)) #w,h
log.plot_epochs(["train_loss", "test_loss"], ax=ax[0], title="Loss_curve");
    ↪#loss_curve
ax[0].legend(loc='upper right', fontsize=10)
log.plot_epochs(["train_acc", "test_acc"], ax=ax[1], title="Acc_curve");
    ↪#acc_curve

```

```

100%|      | 116/116 [00:00<00:00, 300.55it/s]
WARNING:matplotlib.legend:No artists with labels found to put in legend. Note
that artists whose label start with an underscore are ignored when legend() is
called with no argument.
100%|      | 116/116 [00:00<00:00, 297.32it/s]

```



##1.4. Confusion matrix

```

[54]: cm=ConfusionMatrix(task="multiclass", num_classes=10)


```

```

[55]: labels, result=[], []
for data in test_dl:
    _, label=data
    l, a, outputs= net.test_step(data)
    result.append(outputs.argmax(dim=-1))
    labels.append(label.argmax(dim=-1))
cm(torch.flatten(torch.stack(result).to("cpu")), torch.flatten(torch.
    ↪stack(labels).to("cpu")));

```

```
[56]: fig, ax = plt.subplots(figsize=(10, 10))
      cm.plot(ax=ax)
      ax.set_title("Confusion_matrix");
```



```
[ ]: '''Our model is performing well, for an ideal confusion matrix all elements
      ↪ along the diagonal shd be maximum as it means that the true labels were
      ↪ identified correctly'''
```

##1.5. Classification Report

```
[58]: c_report_1=classification_report(torch.flatten(torch.stack(labels).to("cpu")),
      ↪ torch.flatten(torch.stack(result).to("cpu")), output_dict=True)
```

```
df_1= pd.DataFrame(c_report_1)
df_1
```

```
[58]:
```

	0	1	2	3	4 \
precision	0.949555	0.975439	0.962810	0.911597	0.932663
recall	0.980592	0.981465	0.904854	0.951389	0.946939
f1-score	0.964824	0.978443	0.932933	0.931068	0.939747
support	979.000000	1133.000000	1030.000000	1008.000000	980.000000

	5	6	7	8	9 \
precision	0.949292	0.954832	0.963115	0.900293	0.916585
recall	0.904494	0.950837	0.915287	0.946557	0.926587
f1-score	0.926352	0.952830	0.938592	0.922846	0.921559
support	890.000000	956.000000	1027.000000	973.000000	1008.000000

	accuracy	macro avg	weighted avg
precision	0.941506	0.941618	0.942117
recall	0.941506	0.940900	0.941506
f1-score	0.941506	0.940919	0.941474
support	0.941506	9984.000000	9984.000000

#2. Using PyTorch Libraries

##2.1. Building the model

```
[59]: model= nn.Sequential(nn.Flatten(),
                           nn.Linear(28*28, 500), nn.ReLU(),
                           nn.Linear(500, 250), nn.ReLU(),
                           nn.Linear(250, 100), nn.ReLU(),
                           nn.Linear(100, 10), nn.Softmax(dim=1)).to(device)
```

```
[60]: summary(model, input_size=(64, 28*28))
```

```
[60]: =====
=====
Layer (type:depth-idx)                Output Shape                Param #
=====
=====
Sequential                             [64, 10]                    --
  Flatten: 1-1                         [64, 784]                   --
  Linear: 1-2                           [64, 500]                   392,500
  ReLU: 1-3                             [64, 500]                   --
  Linear: 1-4                           [64, 250]                   125,250
  ReLU: 1-5                             [64, 250]                   --
  Linear: 1-6                           [64, 100]                   25,100
  ReLU: 1-7                             [64, 100]                   --
  Linear: 1-8                           [64, 10]                    1,010
  Softmax: 1-9                         [64, 10]                    --
```



```

=====
=====
Total params: 543,860
Trainable params: 543,860
Non-trainable params: 0
Total mult-adds (M): 34.81
=====
=====
Input size (MB): 0.20
Forward/backward pass size (MB): 0.44
Params size (MB): 2.18
Estimated Total Size (MB): 2.82
=====
=====

```

##2.2. Run the model

```
[61]: img, l= next(iter(train_dl))
```

```
[62]: img.shape, l.shape
```

```
[62]: (torch.Size([64, 1, 28, 28]), torch.Size([64, 1, 10]))
```

```
[63]: optimizer=torch.optim.SGD(params= model.parameters(), lr=0.01) #torch.optim.
      ↪ Adam(params= model.parameters()) #Adam converges faster
criterion= cross_entropy
EPOCHS=15
```

```
[64]: def train_step(model, data, criterion, optimizer):
      model.train()
      img, label= data
      img, label= img.to(device), label.to(device).squeeze()
      output= model(img)
      loss= criterion(label, output)
      t_acc= accuracy(label, output)
      optimizer.zero_grad()
      loss.backward()
      optimizer.step()
      return loss.item(), t_acc

def test_step(model, data, criterion):
    model.eval()
    with torch.inference_mode():
        img, label= data
        img, label= img.to(device), label.to(device).squeeze()
        output= model(img)
        loss= criterion(label, output)
```

```

t_acc= accuracy(label, output)
return loss.item(), t_acc

```

```

[65]: log=Report(EPOCHS)
for epoch in range(EPOCHS):
    n=len(train_dl)
    for idx, data in enumerate(train_dl):
        loss, t_acc= train_step(model, data, criterion, optimizer)
        log.record((epoch+(1+idx)/n), train_loss= loss, train_acc= t_acc, end="\r")

    n=len(test_dl)
    for idx, data in enumerate(test_dl):
        loss, t_acc= test_step(model, data, criterion)
        log.record((epoch+(1+idx)/n), test_loss= loss, test_acc= t_acc, end="\r")
    log.report_avgs(epoch+1)

```

```

EPOCH: 1.000 test_loss: 1.739 test_acc: 0.587 train_acc: 0.375 train_loss:
2.185 (30.98s - 433.70s remaining)
EPOCH: 2.000 test_loss: 0.502 test_acc: 0.856 train_acc: 0.781 train_loss:
0.857 (47.11s - 306.18s remaining)
EPOCH: 3.000 test_loss: 0.379 test_acc: 0.888 train_acc: 0.878 train_loss:
0.431 (65.07s - 260.26s remaining)
EPOCH: 4.000 test_loss: 0.330 test_acc: 0.903 train_acc: 0.899 train_loss:
0.356 (81.45s - 223.99s remaining)
EPOCH: 5.000 test_loss: 0.298 test_acc: 0.912 train_acc: 0.910 train_loss:
0.316 (97.29s - 194.58s remaining)
EPOCH: 6.000 test_loss: 0.272 test_acc: 0.920 train_acc: 0.919 train_loss:
0.286 (115.03s - 172.54s remaining)
EPOCH: 7.000 test_loss: 0.248 test_acc: 0.927 train_acc: 0.926 train_loss:
0.260 (131.65s - 150.45s remaining)
EPOCH: 8.000 test_loss: 0.226 test_acc: 0.933 train_acc: 0.933 train_loss:
0.235 (145.55s - 127.35s remaining)
EPOCH: 9.000 test_loss: 0.207 test_acc: 0.939 train_acc: 0.939 train_loss:
0.214 (159.50s - 106.33s remaining)
EPOCH: 10.000 test_loss: 0.190 test_acc: 0.945 train_acc: 0.945 train_loss:
0.195 (171.83s - 85.91s remaining)
EPOCH: 11.000 test_loss: 0.175 test_acc: 0.950 train_acc: 0.949 train_loss:
0.178 (183.91s - 66.88s remaining)
EPOCH: 12.000 test_loss: 0.163 test_acc: 0.953 train_acc: 0.953 train_loss:
0.164 (196.25s - 49.06s remaining)
EPOCH: 13.000 test_loss: 0.152 test_acc: 0.956 train_acc: 0.956 train_loss:
0.152 (208.56s - 32.09s remaining)
EPOCH: 14.000 test_loss: 0.143 test_acc: 0.958 train_acc: 0.960 train_loss:
0.141 (220.67s - 15.76s remaining)
EPOCH: 15.000 test_loss: 0.136 test_acc: 0.959 train_acc: 0.963 train_loss:
0.131 (232.98s - 0.00s remaining)


```

```
[66]: fig, ax= plt.subplots(ncols=2, figsize=(15,5)) #w,h
log.plot_epochs(["train_loss", "test_loss"], ax=ax[0], title="Loss_curve");
    ↪ #loss_curve
ax[0].legend(loc='upper right', fontsize=10)
log.plot_epochs(["train_acc", "test_acc"], ax=ax[1], title="Acc_curve");
    ↪ #acc_curve
```

100%| | 116/116 [00:00<00:00, 301.38it/s]

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

100%| | 116/116 [00:00<00:00, 299.50it/s]




##2.3. Confusion matrix

```
[67]: cm=ConfusionMatrix(task="multiclass", num_classes=10)
```

```
[68]: labels, output=[], []

model.eval()
with torch.inference_mode():
    for data in test_dl:
        im, label=data
        output.append(model(im.to(device)).argmax(dim=-1))
        labels.append(label.argmax(dim=-1))
cm(torch.flatten(torch.stack(output)).to("cpu"), torch.flatten(torch.
    ↪ stack(labels)).to("cpu"));
```

```
[69]: fig, ax = plt.subplots(figsize=(10, 10))
cm.plot(ax=ax)
ax.set_title("Confusion_matrix");
```



```
[ ]: '''Our model is performing well, for an ideal confusion matrix all elements
      ↳ along the diagonal shd be maximum as it means that the true labels were
      ↳ identified correctly'''
```

##2.4. Classification_report

```
[70]: c_report_2=classification_report(torch.flatten(torch.stack(labels).to("cpu")),
      ↳ torch.flatten(torch.stack(result).to("cpu")), output_dict=True)
df_2=pd.DataFrame(c_report_2)
df_2
```

```

[70]:
      0          1          2          3          4  \
precision  0.949555  0.975439  0.962810  0.911597  0.932663
recall     0.980592  0.981465  0.904854  0.951389  0.946939
f1-score   0.964824  0.978443  0.932933  0.931068  0.939747
support    979.000000 1133.000000 1030.000000 1008.000000 980.000000

      5          6          7          8          9  \
precision  0.949292  0.954832  0.963115  0.900293  0.916585
recall     0.904494  0.950837  0.915287  0.946557  0.926587
f1-score   0.926352  0.952830  0.938592  0.922846  0.921559
support    890.000000 956.000000 1027.000000 973.000000 1008.000000

      accuracy  macro avg  weighted avg
precision  0.941506    0.941618    0.942117
recall     0.941506    0.940900    0.941506
f1-score   0.941506    0.940919    0.941474
support    0.941506  9984.000000  9984.000000

```

[70]: