

PA3_na21b033

October 16, 2023

##1.1. Installing necessary packages

```
[1]: !pip install -q torch_snippets
!pip install -q torchinfo
!pip install -q torchmetrics

from torchmetrics import ConfusionMatrix
from torch_snippets import*
from torchinfo import summary
import torch
from torchvision import datasets, transforms
from torch import nn
import matplotlib.pyplot as plt
import sklearn
from sklearn.metrics import classification_report
import pandas as pd
```

```
65.4/65.4 kB
2.3 MB/s eta 0:00:00
115.3/115.3
kB 7.4 MB/s eta 0:00:00
62.5/62.5 kB
7.5 MB/s eta 0:00:00
78.6/78.6 kB
7.8 MB/s eta 0:00:00
Preparing metadata (setup.py) ... done
203.7/203.7
kB 9.4 MB/s eta 0:00:00
4.3/4.3 MB
18.7 MB/s eta 0:00:00
1.6/1.6 MB
31.8 MB/s eta 0:00:00
98.9/98.9 kB
14.0 MB/s eta 0:00:00
3.7/3.7 MB
34.1 MB/s eta 0:00:00
30.6/30.6 MB
```

```

38.8 MB/s eta 0:00:00
169.4/169.4

kB 18.8 MB/s eta 0:00:00
3.2/3.2 MB

46.5 MB/s eta 0:00:00
468.9/468.9

kB 33.8 MB/s eta 0:00:00
Building wheel for typing (setup.py) ... done
805.2/805.2

kB 9.8 MB/s eta 0:00:00

```

```

[2]: device="cuda" if torch.cuda.is_available() else "cpu"
    device #device agnostic code

```

```

[2]: 'cuda'

```

##1.2. Getting the data

```

[3]: transformer= transforms.Compose([transforms.ToTensor()]) #dataset contains imgs
    ↪in PIL format

```

```

[4]: train_ds=datasets.MNIST(root="MNIST/", train=True, download=True, transform=
    ↪transformer)
    #convert labels to one hot encoding
    test_ds=datasets.MNIST(root="MNIST/", train=False, download=True, transform=
    ↪transformer)

    train_dl= torch.utils.data.DataLoader(train_ds, batch_size=32, drop_last=True)
    test_dl= torch.utils.data.DataLoader(test_ds, batch_size=32, drop_last=True)

```

```

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to
MNIST/MNIST/raw/train-images-idx3-ubyte.gz

```

```

100%|          | 9912422/9912422 [00:00<00:00, 142083807.49it/s]

```

```

Extracting MNIST/MNIST/raw/train-images-idx3-ubyte.gz to MNIST/MNIST/raw

```

```

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to
MNIST/MNIST/raw/train-labels-idx1-ubyte.gz

```

```

100%|          | 28881/28881 [00:00<00:00, 115808502.70it/s]

```

Extracting MNIST/MNIST/raw/train-labels-idx1-ubyte.gz to MNIST/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to MNIST/MNIST/raw/t10k-images-idx3-ubyte.gz

100%| | 1648877/1648877 [00:00<00:00, 39557806.99it/s]

Extracting MNIST/MNIST/raw/t10k-images-idx3-ubyte.gz to MNIST/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to MNIST/MNIST/raw/t10k-labels-idx1-ubyte.gz

100%| | 4542/4542 [00:00<00:00, 20843029.29it/s]

Extracting MNIST/MNIST/raw/t10k-labels-idx1-ubyte.gz to MNIST/MNIST/raw

```
[5]: im, label=next(iter(train_dl))
      im.shape, label.shape #labels are not one hot encoded; img_shape-> [28*28]
```

```
[5]: (torch.Size([32, 1, 28, 28]), torch.Size([32]))
```

##1.3. Building naive_rnn

```
[6]: input_size=28
      sequence_length= 28
      hidden_size= 128
      num_layers=2
      num_classes= 10
```

```
[7]: class naive_RNN(nn.Module):
      def __init__(self, input_size, hidden_size, num_layers, num_classes):
          super().__init__()
          self.input_size= input_size
          self.hidden_size= hidden_size
          self.num_layers= num_layers
          self.num_classes= num_classes

          self.rnn_block= nn.RNN(input_size, hidden_size, num_layers,
          ↪batch_first=True)

          self.classification_block= nn.Sequential(nn.Linear(hidden_size,
          ↪num_classes),
                                                    nn.Softmax(dim=-1))
```

```

def forward(self, x):
    h0, c0= torch.zeros(self.num_layers, x.size(0), self.hidden_size).
    ↪to(device), torch.zeros(self.num_layers, x.size(0), self.hidden_size).
    ↪to(device)
    x, h_out= self.rnn_block(x, h0)
    x= self.classification_block(x[:, -1,:])
    return x

def loss_and_accuracy(self, y_true, y_p):
    loss_fn= nn.CrossEntropyLoss()
    loss= loss_fn(y_p, y_true)
    y_p= torch.argmax(y_p, dim=-1)
    acc= torch.sum(torch.eq(y_true, y_p))/len(y_true)
    return loss, acc

```

```

[8]: naive_rnn= naive_RNN(input_size, hidden_size, num_layers, num_classes).
    ↪to(device)
summary(naive_rnn, input_size= (32, 28, 28), col_names= ["input_size",
    ↪"output_size", "num_params"])

```

```

[8]: =====
=====
Layer (type:depth-idx)          Input Shape          Output Shape
Param #
=====
=====
naive_RNN                       [32, 28, 28]         [32, 10]
--
  RNN: 1-1                       [32, 28, 28]         [32, 28, 128]
53,248
  Sequential: 1-2                [32, 128]            [32, 10]
--
    Linear: 2-1                  [32, 128]            [32, 10]
1,290
    Softmax: 2-2                 [32, 10]             [32, 10]
--
=====
=====
Total params: 54,538
Trainable params: 54,538
Non-trainable params: 0
Total mult-adds (M): 47.75
=====
=====
Input size (MB): 0.10
Forward/backward pass size (MB): 0.92

```

Params size (MB): 0.22
Estimated Total Size (MB): 1.24

=====
=====

###1.3.1. Training the model

```
[9]: def train_epoch(model, input, criterion, optimizer, sequence_length=28):
    x, y= input
    x, y= x.to(device).view(-1, sequence_length, input_size), y.to(device)
    model.train()
    output=model(x)
    loss, acc=criterion(y, output)
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
    return loss.item(), acc.item()

def val_epoch(model, input, criterion, sequence_length=28):
    x, y= input
    x, y= x.to(device).view(-1, sequence_length, input_size), y.to(device)
    model.eval()
    with torch.inference_mode():
        output=model(x)
        loss, acc=criterion(y, output)
    return loss.item(), acc.item()
```

```
[10]: EPOCHS=5 #the val acc seems to saturate after 5th epoch
criterion= naive_rnn.loss_and_accuracy
optimizer=torch.optim.Adam(params= naive_rnn.parameters())
```

```
[11]: log= Report(EPOCHS)
for epoch in range(EPOCHS):
    n=len(train_dl)
    for ix, input in enumerate(train_dl):
        train_loss, train_accuracy=train_epoch(naive_rnn, input, criterion,
optimizer)
        log.record(epoch+(ix+1)/n, train_loss=train_loss, train_acc=
train_accuracy, end="\r")
    n=len(test_dl)
    for ix, input in enumerate(test_dl):
        val_loss, val_accuracy=val_epoch(naive_rnn, input, criterion)
        log.record(epoch+(ix+1)/n, val_loss=val_loss, val_acc= val_accuracy,
end="\r")
    log.report_avgs(epoch+1)
```

EPOCH: 1.000 train_loss: 2.012 val_loss: 1.864 train_acc: 0.452 val_acc:

```

0.598 (20.20s - 80.81s remaining)
EPOCH: 2.000 train_loss: 1.905 val_loss: 1.861 train_acc: 0.558 val_acc:
0.603 (34.29s - 51.43s remaining)
EPOCH: 3.000 train_loss: 2.081 val_loss: 1.953 train_acc: 0.375 val_acc:
0.509 (48.35s - 32.24s remaining)
EPOCH: 4.000 train_loss: 2.015 val_loss: 2.252 train_acc: 0.448 val_acc:
0.205 (62.59s - 15.65s remaining)
EPOCH: 5.000 train_loss: 2.135 val_loss: 2.006 train_acc: 0.326 val_acc:
0.454 (76.71s - 0.00s remaining)

```

```

[12]: naive_rnn_results= {"train_acc": train_accuracy, "val_acc": val_accuracy,
    ↪ "train_loss": train_loss, "val_loss": val_loss, "epochs": 5}

```

```

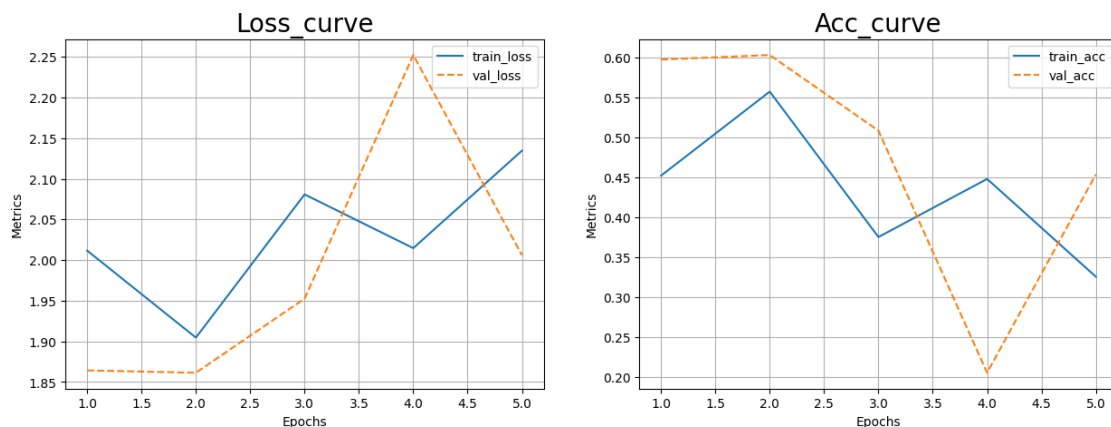
[13]: fig, ax= plt.subplots(ncols=2, figsize=(15,5)) #w,h
    log.plot_epochs(["train_loss", "val_loss"], ax=ax[0], title="Loss_curve");
    ↪ #loss_curve
    ax[0].legend(loc='upper right', fontsize=10)
    log.plot_epochs(["train_acc", "val_acc"], ax=ax[1], title="Acc_curve");
    ↪ #acc_curve

```

```

100%|      | 106/106 [00:00<00:00, 429.38it/s]
WARNING:matplotlib.legend.No artists with labels found to put in legend. Note
that artists whose label start with an underscore are ignored when legend() is
called with no argument.
100%|      | 106/106 [00:00<00:00, 428.96it/s]

```



##1.4. Building LSTM

```

[14]: class LSTM(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers, num_classes):
        super().__init__()
        self.input_size= input_size
        self.hidden_size= hidden_size

```

```

self.num_layers= num_layers
self.num_classes= num_classes

self.lstm_block= nn.LSTM(input_size, hidden_size, num_layers,
↳batch_first=True)

self.classification_block= nn.Sequential(nn.Linear(hidden_size,
↳num_classes),

                                     nn.Softmax(dim=-1))

def forward(self, x):
    h0, c0= torch.zeros(self.num_layers, x.size(0), self.hidden_size).
↳to(device), torch.zeros(self.num_layers, x.size(0), self.hidden_size).
↳to(device)
    x, h_out= self.lstm_block(x, (h0, c0))
    x= self.classification_block(x[:, -1,:])
    return x

def loss_and_accuracy(self, y_true, y_p):
    loss_fn= nn.CrossEntropyLoss()
    loss= loss_fn(y_p, y_true)
    y_p= torch.argmax(y_p, dim=-1)
    acc= torch.sum(torch.eq(y_true, y_p))/len(y_true)
    return loss, acc

```

```

[15]: lstm= LSTM(input_size, hidden_size, num_layers, num_classes).to(device)
summary(lstm, input_size= (32, 28, 28), col_names= ["input_size",
↳"output_size", "num_params"])

```

```

[15]: =====
=====
Layer (type:depth-idx)          Input Shape          Output Shape
Param #
=====
=====
LSTM                            [32, 28, 28]         [32, 10]
--
  LSTM: 1-1                      [32, 28, 28]         [32, 28, 128]
212,992
  Sequential: 1-2                [32, 128]            [32, 10]
--
    Linear: 2-1                  [32, 128]            [32, 10]
1,290
    Softmax: 2-2                 [32, 10]             [32, 10]
--
=====
=====

```

```
Total params: 214,282
Trainable params: 214,282
Non-trainable params: 0
Total mult-adds (M): 190.88
```

```
=====
Input size (MB): 0.10
Forward/backward pass size (MB): 0.92
Params size (MB): 0.86
Estimated Total Size (MB): 1.88
=====
```

###1.4.1. Training the model

```
[16]: EPOCHS=5 #the val acc seems to saturate after 5th epoch
      criterion= lstm.loss_and_accuracy
      optimizer=torch.optim.Adam(params= lstm.parameters())
```

```
[17]: log= Report(EPOCHS)
      for epoch in range(EPOCHS):
          n=len(train_dl)
          for ix, input in enumerate(train_dl):
              train_loss, train_accuracy=train_epoch(lstm, input, criterion, optimizer)
              log.record(epoch+(ix+1)/n, train_loss=train_loss, train_acc=
↳train_accuracy, end="\r")
          n=len(test_dl)
          for ix, input in enumerate(test_dl):
              val_loss, val_accuracy=val_epoch(lstm, input, criterion)
              log.record(epoch+(ix+1)/n, val_loss=val_loss, val_acc= val_accuracy,
↳end="\r")
          log.report_avgs(epoch+1)
```

```
EPOCH: 1.000  train_loss: 1.687  val_loss: 1.562  train_acc: 0.778  val_acc:
0.902  (17.80s - 71.20s remaining)
EPOCH: 2.000  train_loss: 1.522  val_loss: 1.516  train_acc: 0.940  val_acc:
0.945  (33.62s - 50.43s remaining)
EPOCH: 3.000  train_loss: 1.509  val_loss: 1.503  train_acc: 0.952  val_acc:
0.959  (49.39s - 32.93s remaining)
EPOCH: 4.000  train_loss: 1.497  val_loss: 1.506  train_acc: 0.965  val_acc:
0.956  (65.18s - 16.29s remaining)
EPOCH: 5.000  train_loss: 1.495  val_loss: 1.489  train_acc: 0.967  val_acc:
0.972  (81.18s - 0.00s remaining)
```

```
[18]: lstm_results= {"train_acc": train_accuracy, "val_acc": val_accuracy,
↳"train_loss": train_loss, "val_loss": val_loss, "epochs": 5}
```




```
[19]: fig, ax= plt.subplots(ncols=2, figsize=(15,5)) #w,h
log.plot_epochs(["train_loss", "val_loss"], ax=ax[0], title="Loss_curve");
↳#loss_curve
ax[0].legend(loc='upper right', fontsize=10)
log.plot_epochs(["train_acc", "val_acc"], ax=ax[1], title="Acc_curve");
↳#acc_curve
```

100%| | 106/106 [00:00<00:00, 205.34it/s]

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

100%| | 106/106 [00:00<00:00, 220.85it/s]



##1.5. Building GRU

```
[20]: class GRU(nn.Module):
def __init__(self, input_size, hidden_size, num_layers, num_classes):
super().__init__()
self.input_size= input_size
self.hidden_size= hidden_size
self.num_layers= num_layers
self.num_classes= num_classes

self.rnn_block= nn.GRU(input_size, hidden_size, num_layers,
↳batch_first=True)

self.classification_block= nn.Sequential(nn.Linear(hidden_size,
↳num_classes),
nn.Softmax(dim=-1))

def forward(self, x):
```

```

        h0, c0= torch.zeros(self.num_layers, x.size(0), self.hidden_size).
        ↪to(device), torch.zeros(self.num_layers, x.size(0), self.hidden_size).
        ↪to(device)
        x, h_out= self.rnn_block(x,h0)
        x= self.classification_block(x[:, -1,:])
        return x

    def loss_and_accuracy(self, y_true, y_p):
        loss_fn= nn.CrossEntropyLoss()
        loss= loss_fn(y_p, y_true)
        y_p= torch.argmax(y_p, dim=-1)
        acc= torch.sum(torch.eq(y_true, y_p))/len(y_true)
        return loss, acc

```

```

[21]: gru= GRU(input_size, hidden_size, num_layers, num_classes).to(device)
summary(gru, input_size= (32, 28, 28), col_names= ["input_size", "output_size",
        ↪"num_params"])

```

```

[21]: =====
=====
Layer (type:depth-idx)          Input Shape          Output Shape
Param #
=====
=====
GRU                             [32, 28, 28]         [32, 10]
--
  GRU: 1-1                      [32, 28, 28]         [32, 28, 128]
159,744
  Sequential: 1-2               [32, 128]            [32, 10]
--
    Linear: 2-1                 [32, 128]            [32, 10]
1,290
    Softmax: 2-2                [32, 10]             [32, 10]
--
=====
=====
Total params: 161,034
Trainable params: 161,034
Non-trainable params: 0
Total mult-adds (M): 143.17
=====
=====
Input size (MB): 0.10
Forward/backward pass size (MB): 0.92
Params size (MB): 0.64
Estimated Total Size (MB): 1.66
=====

```

=====

###1.5.1. Training the model

```
[22]: EPOCHS=5 #the val acc seems to saturate after 5th epoch
      criterion= gru.loss_and_accuracy
      optimizer=torch.optim.Adam(params= gru.parameters())
```

```
[23]: log= Report(EPOCHS)
      for epoch in range(EPOCHS):
          n=len(train_dl)
          for ix, input in enumerate(train_dl):
              loss, accuracy=train_epoch(gru, input, criterion, optimizer)
              log.record(epoch+(ix+1)/n, train_loss=loss, train_acc= accuracy, end="\r")
          n=len(test_dl)
          for ix, input in enumerate(test_dl):
              loss, accuracy=val_epoch(gru, input, criterion)
              log.record(epoch+(ix+1)/n, val_loss=loss, val_acc= accuracy, end="\r")
          log.report_avgs(epoch+1)
```

```
EPOCH: 1.000  train_loss: 1.637  val_loss: 1.528  train_acc: 0.829  val_acc:
0.934  (14.63s - 58.52s remaining)
EPOCH: 2.000  train_loss: 1.508  val_loss: 1.500  train_acc: 0.955  val_acc:
0.963  (29.07s - 43.60s remaining)
EPOCH: 3.000  train_loss: 1.495  val_loss: 1.491  train_acc: 0.966  val_acc:
0.971  (43.62s - 29.08s remaining)
EPOCH: 4.000  train_loss: 1.489  val_loss: 1.495  train_acc: 0.973  val_acc:
0.967  (58.24s - 14.56s remaining)
EPOCH: 5.000  train_loss: 1.484  val_loss: 1.482  train_acc: 0.977  val_acc:
0.980  (72.74s - 0.00s remaining)
```


```
[24]: gru_results= {"train_acc": train_accuracy, "val_acc": val_accuracy,
                    ↪ "train_loss": train_loss, "val_loss": val_loss, "epochs": 5}
```

```
[25]: fig, ax= plt.subplots(ncols=2, figsize=(15,5)) #w,h
      log.plot_epochs(["train_loss", "val_loss"], ax=ax[0], title="Loss_curve");
      ↪ #loss_curve
      ax[0].legend(loc='upper right', fontsize=10)
      log.plot_epochs(["train_acc", "val_acc"], ax=ax[1], title="Acc_curve");
      ↪ #acc_curve
```

```
100%|      | 106/106 [00:00<00:00, 404.98it/s]
```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

```
100%|      | 106/106 [00:00<00:00, 385.47it/s]
```



##1.6. Building LSTM_bi

```
[26]: class LSTM_bi(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers, num_classes):
        super().__init__()
        self.input_size= input_size
        self.hidden_size= hidden_size
        self.num_layers= num_layers
        self.num_classes= num_classes

        self.lstm_block= nn.LSTM(input_size, hidden_size, num_layers,
        ↪batch_first=True, bidirectional= True)

        self.classification_block= nn.Sequential(nn.Linear(2*hidden_size,
        ↪num_classes),
                                                nn.Softmax(dim=-1))

    def forward(self, x):
        h0, c0= torch.zeros(2*self.num_layers, x.size(0), self.hidden_size).
        ↪to(device), torch.zeros(2*self.num_layers, x.size(0), self.hidden_size).
        ↪to(device)
        x, h_out= self.lstm_block(x, (h0, c0))
        x= self.classification_block(x[:, -1,:])
        return x

    def loss_and_accuracy(self, y_true, y_p):
        loss_fn= nn.CrossEntropyLoss()
        loss= loss_fn(y_p, y_true)
        y_p= torch.argmax(y_p, dim=-1)
        acc= torch.sum(torch.eq(y_true, y_p))/len(y_true)
        return loss, acc
```

```
[27]: lstm_bi= LSTM_bi(input_size, hidden_size, num_layers, num_classes).to(device)
summary(lstm_bi, input_size= (32, 28, 28), col_names= ["input_size", "output_size", "num_params"])
```

```
[27]: =====
=====
Layer (type:depth-idx)          Input Shape          Output Shape
Param #
=====
=====
LSTM_bi                        [32, 28, 28]         [32, 10]
--
  LSTM: 1-1                     [32, 28, 28]         [32, 28, 256]
557,056
  Sequential: 1-2              [32, 256]            [32, 10]
--
    Linear: 2-1                [32, 256]            [32, 10]
2,570
    Softmax: 2-2               [32, 10]             [32, 10]
--
=====
=====
Total params: 559,626
Trainable params: 559,626
Non-trainable params: 0
Total mult-adds (M): 499.20
=====
=====
Input size (MB): 0.10
Forward/backward pass size (MB): 1.84
Params size (MB): 2.24
Estimated Total Size (MB): 4.18
=====
=====
```

###1.6.1. Training the model

```
[28]: EPOCHS=5 #the val acc seems to saturate after 5th epoch
criterion= lstm_bi.loss_and_accuracy
optimizer=torch.optim.Adam(params= lstm_bi.parameters())
```

```
[29]: log= Report(EPOCHS)
for epoch in range(EPOCHS):
    n=len(train_dl)
    for ix, input in enumerate(train_dl):
        train_loss, train_accuracy=train_epoch(lstm_bi, input, criterion, optimizer)
```

```

    log.record(epoch+(ix+1)/n, train_loss=train_loss, train_acc=
↪train_accuracy, end="\r")
    n=len(test_dl)
    for ix, input in enumerate(test_dl):
        val_loss, val_accuracy=val_epoch(lstm_bi, input, criterion)
        log.record(epoch+(ix+1)/n, val_loss=val_loss, val_acc= val_accuracy,
↪end="\r")
    log.report_avgs(epoch+1)

```

```

EPOCH: 1.000  train_loss: 1.683  val_loss: 1.535  train_acc: 0.781  val_acc:
0.929  (18.35s - 73.38s remaining)
EPOCH: 2.000  train_loss: 1.518  val_loss: 1.506  train_acc: 0.945  val_acc:
0.956  (35.55s - 53.33s remaining)
EPOCH: 3.000  train_loss: 1.500  val_loss: 1.495  train_acc: 0.962  val_acc:
0.967  (52.79s - 35.20s remaining)
EPOCH: 4.000  train_loss: 1.496  val_loss: 1.495  train_acc: 0.965  val_acc:
0.967  (70.88s - 17.72s remaining)
EPOCH: 5.000  train_loss: 1.492  val_loss: 1.486  train_acc: 0.970  val_acc:
0.975  (87.94s - 0.00s remaining)

```

```

[30]: lstm_bi_results= {"train_acc": train_accuracy, "val_acc": val_accuracy,
↪"train_loss": train_loss, "val_loss": val_loss, "epochs": 5}

```

```

[31]: fig, ax= plt.subplots(ncols=2, figsize=(15,5)) #w,h
log.plot_epochs(["train_loss", "val_loss"], ax=ax[0], title="Loss_curve");
↪#loss_curve
ax[0].legend(loc='upper right', fontsize=10)
log.plot_epochs(["train_acc", "val_acc"], ax=ax[1], title="Acc_curve");
↪#acc_curve

```

```

100%|      | 106/106 [00:00<00:00, 404.25it/s]


```

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

```

100%|      | 106/106 [00:00<00:00, 410.50it/s]

```



```
[32]: import pandas as pd
data=[naive_rnn_results, lstm_results, gru_results, lstm_bi_results]
index=["naive_rnn", "lstm", "gru", "lstm_bi"]
result_df= pd.DataFrame(data)
result_df.set_index([index], inplace= True)
result_df
```

```
[32]:
```

	train_acc	val_acc	train_loss	val_loss	epochs
naive_rnn	0.50000	0.50000	1.970288	1.938931	5
lstm	1.00000	0.93750	1.461198	1.517358	5
gru	1.00000	0.93750	1.461198	1.517358	5
lstm_bi	0.96875	0.96875	1.498682	1.482065	5

Note that the LSTM model and the BiDirectional LSTM perform on par, there is no gain in additional bidirection. So for the best performing model we shall consider LSTM itself (for lower run time and complexity)


##1.7. Tesing on val_images

```
[33]: ims, labels= next(iter(test_dl))
ims, labels= ims[:10], labels[:10]
ims.shape, labels.shape
```

```
[33]: (torch.Size([10, 1, 28, 28]), torch.Size([10]))
```

```
[36]: with torch.inference_mode():
out= lstm(ims.view(-1, 28,28).to(device))
out= torch.argmax(out, dim=-1)
```

```
[37]: titles=[f"pred: {out[i].item()} true: {labels[i].item()}" for i in_
↪range(len(out))]
subplots(ims.squeeze().numpy(), titles=titles, figsize= (10,5))
```



Our lstm model performs really well.

#2. Adding 2 binary string

##2.1. Custom dataset

```
[38]: def gen_bin(n):
    gen=[]
    for i in range(2**n, 2**(n+1)):
        gen.append(bin(i))
    return gen

def break_bin(num):
    temp= list(num[3:])
    temp= [int(i) for i in temp]
    return temp

def input_prep(s1, s2):
    label, out= [], []
    num1, num2= break_bin(s1), break_bin(s2)
    num1, num2= num1[::-1], num2[::-1]
    sum = bin(int(s1[3:], 2) + int(s2[3:], 2))[2:]
    for idx in range(len(num1)):
        out.append(torch.Tensor([num1[idx], num2[idx]]))
    sum= list(sum)
    sum=sum[::-1]
    if len(sum)<len(num1)+1:
        sum.extend([0]*(len(num1)+1-len(sum)))
    sum= torch.LongTensor([int(i) for i in sum])
    out= torch.vstack(out)
    return out, sum
```



```
[39]: import random

class binary_string_dataset(Dataset):
    def __init__(self, samples, binary_string_len):
        self.samples= samples
        self.string_len= binary_string_len
        self.data, self.target=[], []

        for i in range(samples):
            x= gen_bin(binary_string_len)
            idx_1, idx_2= random.randint(0, len(x)-1), random.randint(0, len(x)-1)
            input, target= input_prep(x[idx_1], x[idx_2])
            self.data.extend([input])
            self.target.extend([target])

    def __len__(self):
        return self.samples

    def __getitem__(self, idx):
        return self.data[idx], self.target[idx]
```

```
[40]: N=10

train_ds= binary_string_dataset(32*10, N)
test_ds= binary_string_dataset(32*3, N)
train_ds[2]
```

```
[40]: (tensor([[0., 1.],
               [1., 0.],
               [0., 1.],
               [1., 1.],
               [0., 0.],
               [0., 1.],
               [0., 1.],
               [1., 0.],
               [1., 0.],
               [1., 0.]]),
       tensor([1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0]))
```

```
[41]: len(train_ds)
```

```
[41]: 320
```

```
[42]: train_dl= torch.utils.data.DataLoader(train_ds, batch_size=32)
       test_dl= torch.utils.data.DataLoader(test_ds, batch_size=32)
```

##2.2. Setting up a LSTM model

```
[43]: input_size=2
```

```
#hyperparameters:
hidden_size=64
sequence_length= N
num_layers=2
num_classes= N
```

```
[44]: class LSTM(nn.Module):
    def __init__(self, input_size, hidden_size, num_layers, num_classes):
        super().__init__()
        self.input_size= input_size
        self.hidden_size= hidden_size
        self.num_layers= num_layers
        self.num_classes= num_classes

        self.lstm_block= nn.LSTM(input_size, hidden_size, num_layers,
        ↪batch_first=True)

        self.classification_block= nn.Sequential(nn.Linear(hidden_size,
        ↪num_classes+1),
                                                    nn.Sigmoid())

    def forward(self, x):
        h0, c0= torch.zeros(self.num_layers, x.size(0), self.hidden_size).
        ↪to(device), torch.zeros(self.num_layers, x.size(0), self.hidden_size).
        ↪to(device)
        x, h_out= self.lstm_block(x, (h0, c0))
        x= self.classification_block(x[:, -1,:])
        return x

    def loss_and_accuracy(self, y_true, y_p):
        loss_fn= nn.BCELoss()
        y_true= y_true.type(torch.float32)
        loss= loss_fn(y_p, y_true)
        y_p= torch.round(y_p)
        acc= torch.sum(torch.eq(y_true, y_p))/(len(y_true)*(N+1))
        return loss, acc
```

```
[45]: lstm= LSTM(input_size, hidden_size, num_layers, num_classes).to(device)
summary(lstm, input_size= (32, 5, 2), col_names= ["input_size", "output_size",
        ↪"num_params"])
```

```
[45]: =====
=====
Layer (type:depth-idx)          Input Shape          Output Shape
Param #
```

```

=====
LSTM                                [32, 5, 2]                [32, 11]
--
  LSTM: 1-1                          [32, 5, 2]                [32, 5, 64]
50,688
  Sequential: 1-2                    [32, 64]                 [32, 11]
--
    Linear: 2-1                      [32, 64]                 [32, 11]
715
    Sigmoid: 2-2                     [32, 11]                 [32, 11]
--
=====
Total params: 51,403
Trainable params: 51,403
Non-trainable params: 0
Total mult-adds (M): 8.13
=====
Input size (MB): 0.00
Forward/backward pass size (MB): 0.08
Params size (MB): 0.21
Estimated Total Size (MB): 0.29
=====
=====

```

```

[46]: EPOCHS=500
      criterion= lstm.loss_and_accuracy
      optimizer=torch.optim.Adam(params= lstm.parameters())

```

##2.3. Running the model

```

[47]: log= Report(EPOCHS)
      for epoch in range(EPOCHS):
          n=len(train_dl)
          for ix, input in enumerate(train_dl):
              train_loss, train_accuracy=train_epoch(lstm, input, criterion, optimizer,
              ↪sequence_length=sequence_length)
              log.record(epoch+(ix+1)/n, train_loss=train_loss, train_acc=
              ↪train_accuracy, end="\r")
          n=len(test_dl)
          for ix, input in enumerate(test_dl):
              val_loss, val_accuracy=val_epoch(lstm, input, criterion,
              ↪sequence_length=sequence_length)
              log.record(epoch+(ix+1)/n, val_loss=val_loss, val_acc= val_accuracy,
              ↪end="\r")

```

```
log.report_avgs(epoch+1)
```

```
EPOCH: 1.000 train_loss: 0.693 val_loss: 0.695 train_acc: 0.505 val_acc:
0.462 (0.05s - 27.06s remaining)
EPOCH: 2.000 train_loss: 0.692 val_loss: 0.694 train_acc: 0.520 val_acc:
0.514 (0.09s - 21.82s remaining)
EPOCH: 3.000 train_loss: 0.691 val_loss: 0.694 train_acc: 0.539 val_acc:
0.506 (0.12s - 19.98s remaining)
EPOCH: 4.000 train_loss: 0.691 val_loss: 0.695 train_acc: 0.530 val_acc:
0.507 (0.15s - 19.01s remaining)
EPOCH: 5.000 train_loss: 0.690 val_loss: 0.695 train_acc: 0.539 val_acc:
0.515 (0.19s - 18.44s remaining)
EPOCH: 6.000 train_loss: 0.690 val_loss: 0.694 train_acc: 0.545 val_acc:
0.514 (0.22s - 18.16s remaining)
EPOCH: 7.000 train_loss: 0.689 val_loss: 0.693 train_acc: 0.548 val_acc:
0.512 (0.25s - 17.94s remaining)
EPOCH: 8.000 train_loss: 0.687 val_loss: 0.693 train_acc: 0.545 val_acc:
0.512 (0.29s - 17.95s remaining)
EPOCH: 9.000 train_loss: 0.685 val_loss: 0.691 train_acc: 0.543 val_acc:
0.513 (0.32s - 17.67s remaining)
EPOCH: 10.000 train_loss: 0.682 val_loss: 0.687 train_acc: 0.544 val_acc:
0.516 (0.36s - 17.41s remaining)
EPOCH: 11.000 train_loss: 0.678 val_loss: 0.684 train_acc: 0.552 val_acc:
0.524 (0.39s - 17.17s remaining)
EPOCH: 12.000 train_loss: 0.675 val_loss: 0.684 train_acc: 0.558 val_acc:
0.520 (0.42s - 17.01s remaining)
EPOCH: 13.000 train_loss: 0.671 val_loss: 0.680 train_acc: 0.553 val_acc:
0.516 (0.45s - 17.01s remaining)
EPOCH: 14.000 train_loss: 0.666 val_loss: 0.671 train_acc: 0.564 val_acc:
0.546 (0.49s - 16.88s remaining)
EPOCH: 15.000 train_loss: 0.661 val_loss: 0.667 train_acc: 0.571 val_acc:
0.557 (0.52s - 16.78s remaining)
EPOCH: 16.000 train_loss: 0.656 val_loss: 0.666 train_acc: 0.575 val_acc:
0.546 (0.56s - 16.89s remaining)
EPOCH: 17.000 train_loss: 0.652 val_loss: 0.666 train_acc: 0.582 val_acc:
0.554 (0.60s - 17.01s remaining)
EPOCH: 18.000 train_loss: 0.649 val_loss: 0.668 train_acc: 0.591 val_acc:
0.545 (0.63s - 16.91s remaining)
EPOCH: 19.000 train_loss: 0.648 val_loss: 0.668 train_acc: 0.590 val_acc:
0.541 (0.67s - 16.94s remaining)
EPOCH: 20.000 train_loss: 0.646 val_loss: 0.665 train_acc: 0.589 val_acc:
0.547 (0.70s - 16.83s remaining)
EPOCH: 21.000 train_loss: 0.642 val_loss: 0.659 train_acc: 0.600 val_acc:
0.550 (0.73s - 16.73s remaining)
EPOCH: 22.000 train_loss: 0.638 val_loss: 0.657 train_acc: 0.606 val_acc:
0.552 (0.76s - 16.62s remaining)
EPOCH: 23.000 train_loss: 0.635 val_loss: 0.658 train_acc: 0.608 val_acc:
```

0.552 (0.80s - 16.56s remaining)
EPOCH: 24.000 train_loss: 0.634 val_loss: 0.658 train_acc: 0.609 val_acc:
0.561 (0.83s - 16.48s remaining)
EPOCH: 25.000 train_loss: 0.630 val_loss: 0.656 train_acc: 0.618 val_acc:
0.559 (0.87s - 16.60s remaining)
EPOCH: 26.000 train_loss: 0.627 val_loss: 0.657 train_acc: 0.622 val_acc:
0.566 (0.91s - 16.51s remaining)
EPOCH: 27.000 train_loss: 0.624 val_loss: 0.658 train_acc: 0.626 val_acc:
0.566 (0.94s - 16.42s remaining)
EPOCH: 28.000 train_loss: 0.622 val_loss: 0.658 train_acc: 0.629 val_acc:
0.562 (0.97s - 16.33s remaining)
EPOCH: 29.000 train_loss: 0.616 val_loss: 0.656 train_acc: 0.640 val_acc:
0.570 (1.00s - 16.26s remaining)
EPOCH: 30.000 train_loss: 0.611 val_loss: 0.655 train_acc: 0.646 val_acc:
0.576 (1.03s - 16.19s remaining)
EPOCH: 31.000 train_loss: 0.607 val_loss: 0.653 train_acc: 0.651 val_acc:
0.579 (1.07s - 16.17s remaining)
EPOCH: 32.000 train_loss: 0.603 val_loss: 0.652 train_acc: 0.660 val_acc:
0.586 (1.11s - 16.18s remaining)
EPOCH: 33.000 train_loss: 0.598 val_loss: 0.648 train_acc: 0.665 val_acc:
0.589 (1.14s - 16.12s remaining)
EPOCH: 34.000 train_loss: 0.592 val_loss: 0.642 train_acc: 0.674 val_acc:
0.600 (1.17s - 16.06s remaining)
EPOCH: 35.000 train_loss: 0.582 val_loss: 0.638 train_acc: 0.691 val_acc:
0.598 (1.20s - 15.99s remaining)
EPOCH: 36.000 train_loss: 0.572 val_loss: 0.633 train_acc: 0.704 val_acc:
0.613 (1.24s - 15.93s remaining)
EPOCH: 37.000 train_loss: 0.561 val_loss: 0.626 train_acc: 0.714 val_acc:
0.622 (1.28s - 15.99s remaining)
EPOCH: 38.000 train_loss: 0.549 val_loss: 0.613 train_acc: 0.728 val_acc:
0.633 (1.31s - 15.96s remaining)
EPOCH: 39.000 train_loss: 0.535 val_loss: 0.599 train_acc: 0.742 val_acc:
0.658 (1.35s - 15.94s remaining)
EPOCH: 40.000 train_loss: 0.522 val_loss: 0.584 train_acc: 0.748 val_acc:
0.669 (1.38s - 15.90s remaining)
EPOCH: 41.000 train_loss: 0.511 val_loss: 0.574 train_acc: 0.752 val_acc:
0.680 (1.42s - 15.86s remaining)
EPOCH: 42.000 train_loss: 0.502 val_loss: 0.565 train_acc: 0.761 val_acc:
0.683 (1.45s - 15.82s remaining)
EPOCH: 43.000 train_loss: 0.490 val_loss: 0.555 train_acc: 0.776 val_acc:
0.695 (1.49s - 15.82s remaining)
EPOCH: 44.000 train_loss: 0.479 val_loss: 0.544 train_acc: 0.786 val_acc:
0.705 (1.52s - 15.77s remaining)
EPOCH: 45.000 train_loss: 0.468 val_loss: 0.535 train_acc: 0.796 val_acc:
0.722 (1.55s - 15.71s remaining)
EPOCH: 46.000 train_loss: 0.456 val_loss: 0.523 train_acc: 0.801 val_acc:
0.731 (1.59s - 15.66s remaining)
EPOCH: 47.000 train_loss: 0.445 val_loss: 0.513 train_acc: 0.814 val_acc:

0.752 (1.62s - 15.66s remaining)
EPOCH: 48.000 train_loss: 0.434 val_loss: 0.503 train_acc: 0.821 val_acc:
0.759 (1.66s - 15.62s remaining)
EPOCH: 49.000 train_loss: 0.425 val_loss: 0.496 train_acc: 0.827 val_acc:
0.757 (1.70s - 15.61s remaining)
EPOCH: 50.000 train_loss: 0.414 val_loss: 0.488 train_acc: 0.834 val_acc:
0.765 (1.73s - 15.57s remaining)
EPOCH: 51.000 train_loss: 0.407 val_loss: 0.483 train_acc: 0.837 val_acc:
0.765 (1.76s - 15.50s remaining)
EPOCH: 52.000 train_loss: 0.399 val_loss: 0.480 train_acc: 0.840 val_acc:
0.759 (1.79s - 15.46s remaining)
EPOCH: 53.000 train_loss: 0.394 val_loss: 0.470 train_acc: 0.842 val_acc:
0.766 (1.83s - 15.41s remaining)
EPOCH: 54.000 train_loss: 0.386 val_loss: 0.466 train_acc: 0.843 val_acc:
0.775 (1.86s - 15.36s remaining)
EPOCH: 55.000 train_loss: 0.385 val_loss: 0.465 train_acc: 0.843 val_acc:
0.777 (1.90s - 15.34s remaining)
EPOCH: 56.000 train_loss: 0.378 val_loss: 0.453 train_acc: 0.851 val_acc:
0.791 (1.93s - 15.30s remaining)
EPOCH: 57.000 train_loss: 0.371 val_loss: 0.442 train_acc: 0.852 val_acc:
0.795 (1.96s - 15.25s remaining)
EPOCH: 58.000 train_loss: 0.364 val_loss: 0.427 train_acc: 0.859 val_acc:
0.799 (2.00s - 15.25s remaining)
EPOCH: 59.000 train_loss: 0.358 val_loss: 0.412 train_acc: 0.857 val_acc:
0.814 (2.03s - 15.20s remaining)
EPOCH: 60.000 train_loss: 0.352 val_loss: 0.406 train_acc: 0.861 val_acc:
0.818 (2.07s - 15.15s remaining)
EPOCH: 61.000 train_loss: 0.341 val_loss: 0.392 train_acc: 0.867 val_acc:
0.836 (2.10s - 15.13s remaining)
EPOCH: 62.000 train_loss: 0.323 val_loss: 0.385 train_acc: 0.879 val_acc:
0.838 (2.13s - 15.08s remaining)
EPOCH: 63.000 train_loss: 0.316 val_loss: 0.376 train_acc: 0.883 val_acc:
0.834 (2.17s - 15.03s remaining)
EPOCH: 64.000 train_loss: 0.309 val_loss: 0.367 train_acc: 0.887 val_acc:
0.843 (2.20s - 14.99s remaining)
EPOCH: 65.000 train_loss: 0.300 val_loss: 0.361 train_acc: 0.896 val_acc:
0.850 (2.24s - 14.98s remaining)
EPOCH: 66.000 train_loss: 0.293 val_loss: 0.355 train_acc: 0.900 val_acc:
0.848 (2.27s - 14.93s remaining)
EPOCH: 67.000 train_loss: 0.286 val_loss: 0.349 train_acc: 0.903 val_acc:
0.850 (2.31s - 14.91s remaining)
EPOCH: 68.000 train_loss: 0.278 val_loss: 0.343 train_acc: 0.909 val_acc:
0.854 (2.34s - 14.85s remaining)
EPOCH: 69.000 train_loss: 0.271 val_loss: 0.337 train_acc: 0.911 val_acc:
0.862 (2.37s - 14.80s remaining)
EPOCH: 70.000 train_loss: 0.264 val_loss: 0.332 train_acc: 0.916 val_acc:
0.868 (2.40s - 14.75s remaining)
EPOCH: 71.000 train_loss: 0.257 val_loss: 0.327 train_acc: 0.921 val_acc:

0.871 (2.43s - 14.70s remaining)
EPOCH: 72.000 train_loss: 0.250 val_loss: 0.323 train_acc: 0.922 val_acc:
0.875 (2.46s - 14.65s remaining)
EPOCH: 73.000 train_loss: 0.244 val_loss: 0.320 train_acc: 0.925 val_acc:
0.876 (2.50s - 14.62s remaining)
EPOCH: 74.000 train_loss: 0.238 val_loss: 0.316 train_acc: 0.929 val_acc:
0.876 (2.53s - 14.58s remaining)
EPOCH: 75.000 train_loss: 0.232 val_loss: 0.313 train_acc: 0.931 val_acc:
0.877 (2.56s - 14.53s remaining)
EPOCH: 76.000 train_loss: 0.227 val_loss: 0.307 train_acc: 0.934 val_acc:
0.881 (2.60s - 14.53s remaining)
EPOCH: 77.000 train_loss: 0.223 val_loss: 0.302 train_acc: 0.936 val_acc:
0.887 (2.65s - 14.55s remaining)
EPOCH: 78.000 train_loss: 0.219 val_loss: 0.301 train_acc: 0.938 val_acc:
0.883 (2.68s - 14.52s remaining)
EPOCH: 79.000 train_loss: 0.217 val_loss: 0.302 train_acc: 0.938 val_acc:
0.879 (2.72s - 14.50s remaining)
EPOCH: 80.000 train_loss: 0.216 val_loss: 0.296 train_acc: 0.941 val_acc:
0.885 (2.75s - 14.46s remaining)
EPOCH: 81.000 train_loss: 0.211 val_loss: 0.290 train_acc: 0.942 val_acc:
0.887 (2.79s - 14.41s remaining)
EPOCH: 82.000 train_loss: 0.204 val_loss: 0.278 train_acc: 0.947 val_acc:
0.896 (2.83s - 14.41s remaining)
EPOCH: 83.000 train_loss: 0.191 val_loss: 0.271 train_acc: 0.957 val_acc:
0.897 (2.86s - 14.37s remaining)
EPOCH: 84.000 train_loss: 0.184 val_loss: 0.265 train_acc: 0.963 val_acc:
0.908 (2.89s - 14.33s remaining)
EPOCH: 85.000 train_loss: 0.176 val_loss: 0.261 train_acc: 0.967 val_acc:
0.907 (2.93s - 14.32s remaining)
EPOCH: 86.000 train_loss: 0.168 val_loss: 0.256 train_acc: 0.971 val_acc:
0.914 (2.97s - 14.28s remaining)
EPOCH: 87.000 train_loss: 0.162 val_loss: 0.252 train_acc: 0.974 val_acc:
0.912 (3.01s - 14.29s remaining)
EPOCH: 88.000 train_loss: 0.157 val_loss: 0.249 train_acc: 0.975 val_acc:
0.911 (3.04s - 14.25s remaining)
EPOCH: 89.000 train_loss: 0.153 val_loss: 0.245 train_acc: 0.977 val_acc:
0.912 (3.08s - 14.20s remaining)
EPOCH: 90.000 train_loss: 0.149 val_loss: 0.241 train_acc: 0.978 val_acc:
0.915 (3.11s - 14.17s remaining)
EPOCH: 91.000 train_loss: 0.144 val_loss: 0.236 train_acc: 0.980 val_acc:
0.920 (3.14s - 14.13s remaining)
EPOCH: 92.000 train_loss: 0.140 val_loss: 0.232 train_acc: 0.981 val_acc:
0.925 (3.18s - 14.09s remaining)
EPOCH: 93.000 train_loss: 0.136 val_loss: 0.228 train_acc: 0.982 val_acc:
0.927 (3.22s - 14.07s remaining)
EPOCH: 94.000 train_loss: 0.132 val_loss: 0.226 train_acc: 0.983 val_acc:
0.929 (3.25s - 14.05s remaining)
EPOCH: 95.000 train_loss: 0.128 val_loss: 0.223 train_acc: 0.986 val_acc:

0.928 (3.29s - 14.01s remaining)
EPOCH: 96.000 train_loss: 0.124 val_loss: 0.221 train_acc: 0.989 val_acc:
0.927 (3.32s - 13.99s remaining)
EPOCH: 97.000 train_loss: 0.120 val_loss: 0.220 train_acc: 0.989 val_acc:
0.930 (3.36s - 13.94s remaining)
EPOCH: 98.000 train_loss: 0.117 val_loss: 0.220 train_acc: 0.990 val_acc:
0.928 (3.39s - 13.90s remaining)
EPOCH: 99.000 train_loss: 0.115 val_loss: 0.222 train_acc: 0.991 val_acc:
0.929 (3.42s - 13.86s remaining)
EPOCH: 100.000 train_loss: 0.115 val_loss: 0.221 train_acc: 0.990 val_acc:
0.929 (3.45s - 13.82s remaining)
EPOCH: 101.000 train_loss: 0.118 val_loss: 0.211 train_acc: 0.989 val_acc:
0.933 (3.49s - 13.78s remaining)
EPOCH: 102.000 train_loss: 0.116 val_loss: 0.212 train_acc: 0.989 val_acc:
0.933 (3.52s - 13.74s remaining)
EPOCH: 103.000 train_loss: 0.117 val_loss: 0.198 train_acc: 0.982 val_acc:
0.944 (3.56s - 13.71s remaining)
EPOCH: 104.000 train_loss: 0.107 val_loss: 0.190 train_acc: 0.986 val_acc:
0.952 (3.59s - 13.67s remaining)
EPOCH: 105.000 train_loss: 0.101 val_loss: 0.188 train_acc: 0.991 val_acc:
0.945 (3.62s - 13.63s remaining)
EPOCH: 106.000 train_loss: 0.097 val_loss: 0.181 train_acc: 0.994 val_acc:
0.947 (3.66s - 13.61s remaining)
EPOCH: 107.000 train_loss: 0.092 val_loss: 0.179 train_acc: 0.994 val_acc:
0.952 (3.69s - 13.57s remaining)
EPOCH: 108.000 train_loss: 0.089 val_loss: 0.176 train_acc: 0.996 val_acc:
0.950 (3.73s - 13.54s remaining)
EPOCH: 109.000 train_loss: 0.086 val_loss: 0.175 train_acc: 0.995 val_acc:
0.951 (3.76s - 13.50s remaining)
EPOCH: 110.000 train_loss: 0.083 val_loss: 0.174 train_acc: 0.997 val_acc:
0.950 (3.80s - 13.46s remaining)
EPOCH: 111.000 train_loss: 0.081 val_loss: 0.172 train_acc: 0.997 val_acc:
0.953 (3.83s - 13.42s remaining)
EPOCH: 112.000 train_loss: 0.079 val_loss: 0.170 train_acc: 0.997 val_acc:
0.954 (3.86s - 13.38s remaining)
EPOCH: 113.000 train_loss: 0.077 val_loss: 0.167 train_acc: 0.998 val_acc:
0.954 (3.90s - 13.34s remaining)
EPOCH: 114.000 train_loss: 0.075 val_loss: 0.164 train_acc: 0.999 val_acc:
0.955 (3.93s - 13.30s remaining)
EPOCH: 115.000 train_loss: 0.073 val_loss: 0.162 train_acc: 0.998 val_acc:
0.955 (3.97s - 13.29s remaining)
EPOCH: 116.000 train_loss: 0.071 val_loss: 0.161 train_acc: 0.999 val_acc:
0.955 (4.00s - 13.25s remaining)
EPOCH: 117.000 train_loss: 0.069 val_loss: 0.160 train_acc: 0.999 val_acc:
0.956 (4.04s - 13.22s remaining)
EPOCH: 118.000 train_loss: 0.067 val_loss: 0.159 train_acc: 1.000 val_acc:
0.957 (4.07s - 13.18s remaining)
EPOCH: 119.000 train_loss: 0.065 val_loss: 0.158 train_acc: 1.000 val_acc:

0.957	(4.11s - 13.15s remaining)				
EPOCH: 120.000	train_loss: 0.063	val_loss: 0.157	train_acc: 1.000	val_acc:	
0.955	(4.14s - 13.12s remaining)				
EPOCH: 121.000	train_loss: 0.062	val_loss: 0.156	train_acc: 1.000	val_acc:	
0.956	(4.18s - 13.09s remaining)				
EPOCH: 122.000	train_loss: 0.060	val_loss: 0.154	train_acc: 1.000	val_acc:	
0.957	(4.21s - 13.05s remaining)				
EPOCH: 123.000	train_loss: 0.059	val_loss: 0.153	train_acc: 1.000	val_acc:	
0.958	(4.24s - 13.01s remaining)				
EPOCH: 124.000	train_loss: 0.057	val_loss: 0.151	train_acc: 1.000	val_acc:	
0.958	(4.28s - 12.97s remaining)				
EPOCH: 125.000	train_loss: 0.056	val_loss: 0.149	train_acc: 1.000	val_acc:	
0.960	(4.31s - 12.93s remaining)				
EPOCH: 126.000	train_loss: 0.055	val_loss: 0.148	train_acc: 1.000	val_acc:	
0.960	(4.34s - 12.89s remaining)				
EPOCH: 127.000	train_loss: 0.053	val_loss: 0.146	train_acc: 1.000	val_acc:	
0.960	(4.38s - 12.86s remaining)				
EPOCH: 128.000	train_loss: 0.052	val_loss: 0.145	train_acc: 1.000	val_acc:	
0.960	(4.42s - 12.84s remaining)				
EPOCH: 129.000	train_loss: 0.051	val_loss: 0.145	train_acc: 1.000	val_acc:	
0.961	(4.45s - 12.81s remaining)				
EPOCH: 130.000	train_loss: 0.050	val_loss: 0.144	train_acc: 1.000	val_acc:	
0.960	(4.49s - 12.77s remaining)				
EPOCH: 131.000	train_loss: 0.048	val_loss: 0.143	train_acc: 1.000	val_acc:	
0.959	(4.52s - 12.73s remaining)				
EPOCH: 132.000	train_loss: 0.047	val_loss: 0.143	train_acc: 1.000	val_acc:	
0.961	(4.56s - 12.70s remaining)				
EPOCH: 133.000	train_loss: 0.046	val_loss: 0.141	train_acc: 1.000	val_acc:	
0.962	(4.59s - 12.67s remaining)				
EPOCH: 134.000	train_loss: 0.045	val_loss: 0.140	train_acc: 1.000	val_acc:	
0.963	(4.63s - 12.64s remaining)				
EPOCH: 135.000	train_loss: 0.044	val_loss: 0.139	train_acc: 1.000	val_acc:	
0.965	(4.67s - 12.61s remaining)				
EPOCH: 136.000	train_loss: 0.043	val_loss: 0.137	train_acc: 1.000	val_acc:	
0.965	(4.70s - 12.57s remaining)				
EPOCH: 137.000	train_loss: 0.042	val_loss: 0.136	train_acc: 1.000	val_acc:	
0.966	(4.73s - 12.53s remaining)				
EPOCH: 138.000	train_loss: 0.041	val_loss: 0.134	train_acc: 1.000	val_acc:	
0.967	(4.77s - 12.50s remaining)				
EPOCH: 139.000	train_loss: 0.040	val_loss: 0.133	train_acc: 1.000	val_acc:	
0.966	(4.80s - 12.46s remaining)				
EPOCH: 140.000	train_loss: 0.039	val_loss: 0.132	train_acc: 1.000	val_acc:	
0.966	(4.83s - 12.43s remaining)				
EPOCH: 141.000	train_loss: 0.039	val_loss: 0.131	train_acc: 1.000	val_acc:	
0.965	(4.87s - 12.39s remaining)				
EPOCH: 142.000	train_loss: 0.038	val_loss: 0.130	train_acc: 1.000	val_acc:	
0.967	(4.90s - 12.35s remaining)				
EPOCH: 143.000	train_loss: 0.037	val_loss: 0.129	train_acc: 1.000	val_acc:	

```

0.966 (4.93s - 12.31s remaining)
EPOCH: 144.000 train_loss: 0.036 val_loss: 0.128 train_acc: 1.000 val_acc:
0.967 (4.96s - 12.27s remaining)
EPOCH: 145.000 train_loss: 0.035 val_loss: 0.128 train_acc: 1.000 val_acc:
0.967 (5.00s - 12.25s remaining)
EPOCH: 146.000 train_loss: 0.035 val_loss: 0.127 train_acc: 1.000 val_acc:
0.970 (5.04s - 12.21s remaining)
EPOCH: 147.000 train_loss: 0.034 val_loss: 0.127 train_acc: 1.000 val_acc:
0.971 (5.07s - 12.18s remaining)
EPOCH: 148.000 train_loss: 0.033 val_loss: 0.126 train_acc: 1.000 val_acc:
0.972 (5.10s - 12.14s remaining)
EPOCH: 149.000 train_loss: 0.033 val_loss: 0.126 train_acc: 1.000 val_acc:
0.973 (5.14s - 12.11s remaining)
EPOCH: 150.000 train_loss: 0.032 val_loss: 0.125 train_acc: 1.000 val_acc:
0.972 (5.18s - 12.08s remaining)
EPOCH: 151.000 train_loss: 0.032 val_loss: 0.124 train_acc: 1.000 val_acc:
0.973 (5.21s - 12.05s remaining)
EPOCH: 152.000 train_loss: 0.031 val_loss: 0.122 train_acc: 1.000 val_acc:
0.973 (5.24s - 12.00s remaining)
EPOCH: 153.000 train_loss: 0.031 val_loss: 0.120 train_acc: 1.000 val_acc:
0.971 (5.28s - 11.97s remaining)
EPOCH: 154.000 train_loss: 0.031 val_loss: 0.118 train_acc: 1.000 val_acc:
0.967 (5.31s - 11.93s remaining)
EPOCH: 155.000 train_loss: 0.030 val_loss: 0.118 train_acc: 1.000 val_acc:
0.966 (5.34s - 11.90s remaining)
EPOCH: 156.000 train_loss: 0.030 val_loss: 0.117 train_acc: 1.000 val_acc:
0.967 (5.38s - 11.86s remaining)
EPOCH: 157.000 train_loss: 0.029 val_loss: 0.115 train_acc: 1.000 val_acc:
0.970 (5.42s - 11.84s remaining)
EPOCH: 158.000 train_loss: 0.029 val_loss: 0.113 train_acc: 1.000 val_acc:
0.971 (5.45s - 11.80s remaining)
EPOCH: 159.000 train_loss: 0.028 val_loss: 0.115 train_acc: 1.000 val_acc:
0.973 (5.49s - 11.77s remaining)
EPOCH: 160.000 train_loss: 0.028 val_loss: 0.118 train_acc: 1.000 val_acc:
0.973 (5.52s - 11.73s remaining)
EPOCH: 161.000 train_loss: 0.028 val_loss: 0.119 train_acc: 1.000 val_acc:
0.973 (5.56s - 11.70s remaining)
EPOCH: 162.000 train_loss: 0.027 val_loss: 0.115 train_acc: 1.000 val_acc:
0.973 (5.59s - 11.67s remaining)
EPOCH: 163.000 train_loss: 0.026 val_loss: 0.111 train_acc: 1.000 val_acc:
0.969 (5.63s - 11.64s remaining)
EPOCH: 164.000 train_loss: 0.026 val_loss: 0.107 train_acc: 1.000 val_acc:
0.968 (5.67s - 11.61s remaining)
EPOCH: 165.000 train_loss: 0.025 val_loss: 0.106 train_acc: 1.000 val_acc:
0.973 (5.71s - 11.59s remaining)
EPOCH: 166.000 train_loss: 0.024 val_loss: 0.110 train_acc: 1.000 val_acc:
0.973 (5.74s - 11.56s remaining)
EPOCH: 167.000 train_loss: 0.024 val_loss: 0.111 train_acc: 1.000 val_acc:

```

0.974	(5.78s - 11.52s remaining)				
EPOCH: 168.000	train_loss: 0.023	val_loss: 0.108	train_acc: 1.000	val_acc:	
0.975	(5.81s - 11.49s remaining)				
EPOCH: 169.000	train_loss: 0.023	val_loss: 0.107	train_acc: 1.000	val_acc:	
0.973	(5.85s - 11.45s remaining)				
EPOCH: 170.000	train_loss: 0.023	val_loss: 0.105	train_acc: 1.000	val_acc:	
0.973	(5.88s - 11.42s remaining)				
EPOCH: 171.000	train_loss: 0.022	val_loss: 0.105	train_acc: 1.000	val_acc:	
0.974	(5.91s - 11.38s remaining)				
EPOCH: 172.000	train_loss: 0.022	val_loss: 0.106	train_acc: 1.000	val_acc:	
0.975	(5.95s - 11.34s remaining)				
EPOCH: 173.000	train_loss: 0.021	val_loss: 0.106	train_acc: 1.000	val_acc:	
0.975	(5.98s - 11.31s remaining)				
EPOCH: 174.000	train_loss: 0.021	val_loss: 0.105	train_acc: 1.000	val_acc:	
0.975	(6.04s - 11.31s remaining)				
EPOCH: 175.000	train_loss: 0.021	val_loss: 0.103	train_acc: 1.000	val_acc:	
0.976	(6.07s - 11.27s remaining)				
EPOCH: 176.000	train_loss: 0.020	val_loss: 0.103	train_acc: 1.000	val_acc:	
0.975	(6.10s - 11.23s remaining)				
EPOCH: 177.000	train_loss: 0.020	val_loss: 0.102	train_acc: 1.000	val_acc:	
0.975	(6.14s - 11.20s remaining)				
EPOCH: 178.000	train_loss: 0.020	val_loss: 0.102	train_acc: 1.000	val_acc:	
0.974	(6.17s - 11.16s remaining)				
EPOCH: 179.000	train_loss: 0.019	val_loss: 0.102	train_acc: 1.000	val_acc:	
0.975	(6.20s - 11.12s remaining)				
EPOCH: 180.000	train_loss: 0.019	val_loss: 0.102	train_acc: 1.000	val_acc:	
0.975	(6.24s - 11.09s remaining)				
EPOCH: 181.000	train_loss: 0.019	val_loss: 0.101	train_acc: 1.000	val_acc:	
0.974	(6.27s - 11.06s remaining)				
EPOCH: 182.000	train_loss: 0.019	val_loss: 0.100	train_acc: 1.000	val_acc:	
0.974	(6.31s - 11.02s remaining)				
EPOCH: 183.000	train_loss: 0.018	val_loss: 0.100	train_acc: 1.000	val_acc:	
0.975	(6.34s - 10.98s remaining)				
EPOCH: 184.000	train_loss: 0.018	val_loss: 0.100	train_acc: 1.000	val_acc:	
0.974	(6.37s - 10.94s remaining)				
EPOCH: 185.000	train_loss: 0.018	val_loss: 0.099	train_acc: 1.000	val_acc:	
0.974	(6.40s - 10.90s remaining)				
EPOCH: 186.000	train_loss: 0.017	val_loss: 0.099	train_acc: 1.000	val_acc:	
0.975	(6.44s - 10.88s remaining)				
EPOCH: 187.000	train_loss: 0.017	val_loss: 0.099	train_acc: 1.000	val_acc:	
0.976	(6.47s - 10.84s remaining)				
EPOCH: 188.000	train_loss: 0.017	val_loss: 0.099	train_acc: 1.000	val_acc:	
0.974	(6.51s - 10.80s remaining)				
EPOCH: 189.000	train_loss: 0.017	val_loss: 0.098	train_acc: 1.000	val_acc:	
0.974	(6.54s - 10.76s remaining)				
EPOCH: 190.000	train_loss: 0.017	val_loss: 0.097	train_acc: 1.000	val_acc:	
0.974	(6.57s - 10.72s remaining)				
EPOCH: 191.000	train_loss: 0.016	val_loss: 0.097	train_acc: 1.000	val_acc:	

```

0.974 (6.60s - 10.68s remaining)
EPOCH: 192.000 train_loss: 0.016 val_loss: 0.097 train_acc: 1.000 val_acc:
0.974 (6.64s - 10.65s remaining)
EPOCH: 193.000 train_loss: 0.016 val_loss: 0.096 train_acc: 1.000 val_acc:
0.974 (6.67s - 10.61s remaining)
EPOCH: 194.000 train_loss: 0.016 val_loss: 0.096 train_acc: 1.000 val_acc:
0.975 (6.71s - 10.59s remaining)
EPOCH: 195.000 train_loss: 0.015 val_loss: 0.096 train_acc: 1.000 val_acc:
0.975 (6.74s - 10.55s remaining)
EPOCH: 196.000 train_loss: 0.015 val_loss: 0.096 train_acc: 1.000 val_acc:
0.976 (6.78s - 10.51s remaining)
EPOCH: 197.000 train_loss: 0.015 val_loss: 0.096 train_acc: 1.000 val_acc:
0.976 (6.81s - 10.47s remaining)
EPOCH: 198.000 train_loss: 0.015 val_loss: 0.095 train_acc: 1.000 val_acc:
0.976 (6.85s - 10.44s remaining)
EPOCH: 199.000 train_loss: 0.015 val_loss: 0.094 train_acc: 1.000 val_acc:
0.975 (6.88s - 10.41s remaining)
EPOCH: 200.000 train_loss: 0.014 val_loss: 0.094 train_acc: 1.000 val_acc:
0.975 (6.91s - 10.37s remaining)
EPOCH: 201.000 train_loss: 0.014 val_loss: 0.094 train_acc: 1.000 val_acc:
0.975 (6.95s - 10.33s remaining)
EPOCH: 202.000 train_loss: 0.014 val_loss: 0.094 train_acc: 1.000 val_acc:
0.976 (6.98s - 10.29s remaining)
EPOCH: 203.000 train_loss: 0.014 val_loss: 0.093 train_acc: 1.000 val_acc:
0.976 (7.01s - 10.26s remaining)
EPOCH: 204.000 train_loss: 0.014 val_loss: 0.093 train_acc: 1.000 val_acc:
0.976 (7.05s - 10.23s remaining)
EPOCH: 205.000 train_loss: 0.014 val_loss: 0.093 train_acc: 1.000 val_acc:
0.976 (7.09s - 10.20s remaining)
EPOCH: 206.000 train_loss: 0.013 val_loss: 0.093 train_acc: 1.000 val_acc:
0.976 (7.12s - 10.16s remaining)
EPOCH: 207.000 train_loss: 0.013 val_loss: 0.093 train_acc: 1.000 val_acc:
0.976 (7.15s - 10.13s remaining)
EPOCH: 208.000 train_loss: 0.013 val_loss: 0.092 train_acc: 1.000 val_acc:
0.976 (7.19s - 10.09s remaining)
EPOCH: 209.000 train_loss: 0.013 val_loss: 0.092 train_acc: 1.000 val_acc:
0.976 (7.22s - 10.05s remaining)
EPOCH: 210.000 train_loss: 0.013 val_loss: 0.091 train_acc: 1.000 val_acc:
0.976 (7.26s - 10.03s remaining)
EPOCH: 211.000 train_loss: 0.013 val_loss: 0.091 train_acc: 1.000 val_acc:
0.976 (7.29s - 9.99s remaining)
EPOCH: 212.000 train_loss: 0.012 val_loss: 0.091 train_acc: 1.000 val_acc:
0.976 (7.33s - 9.95s remaining)
EPOCH: 213.000 train_loss: 0.012 val_loss: 0.091 train_acc: 1.000 val_acc:
0.976 (7.36s - 9.91s remaining)
EPOCH: 214.000 train_loss: 0.012 val_loss: 0.090 train_acc: 1.000 val_acc:
0.976 (7.39s - 9.88s remaining)
EPOCH: 215.000 train_loss: 0.012 val_loss: 0.091 train_acc: 1.000 val_acc:

```

0.976 (7.42s - 9.84s remaining)
EPOCH: 216.000 train_loss: 0.012 val_loss: 0.091 train_acc: 1.000 val_acc:
0.976 (7.46s - 9.81s remaining)
EPOCH: 217.000 train_loss: 0.012 val_loss: 0.091 train_acc: 1.000 val_acc:
0.976 (7.50s - 9.78s remaining)
EPOCH: 218.000 train_loss: 0.012 val_loss: 0.090 train_acc: 1.000 val_acc:
0.977 (7.54s - 9.75s remaining)
EPOCH: 219.000 train_loss: 0.011 val_loss: 0.090 train_acc: 1.000 val_acc:
0.977 (7.57s - 9.71s remaining)
EPOCH: 220.000 train_loss: 0.011 val_loss: 0.089 train_acc: 1.000 val_acc:
0.977 (7.60s - 9.67s remaining)
EPOCH: 221.000 train_loss: 0.011 val_loss: 0.089 train_acc: 1.000 val_acc:
0.977 (7.63s - 9.63s remaining)
EPOCH: 222.000 train_loss: 0.011 val_loss: 0.088 train_acc: 1.000 val_acc:
0.976 (7.67s - 9.60s remaining)
EPOCH: 223.000 train_loss: 0.011 val_loss: 0.088 train_acc: 1.000 val_acc:
0.977 (7.70s - 9.57s remaining)
EPOCH: 224.000 train_loss: 0.011 val_loss: 0.088 train_acc: 1.000 val_acc:
0.977 (7.74s - 9.54s remaining)
EPOCH: 225.000 train_loss: 0.011 val_loss: 0.087 train_acc: 1.000 val_acc:
0.977 (7.77s - 9.50s remaining)
EPOCH: 226.000 train_loss: 0.011 val_loss: 0.087 train_acc: 1.000 val_acc:
0.976 (7.81s - 9.46s remaining)
EPOCH: 227.000 train_loss: 0.011 val_loss: 0.087 train_acc: 1.000 val_acc:
0.976 (7.84s - 9.43s remaining)
EPOCH: 228.000 train_loss: 0.010 val_loss: 0.088 train_acc: 1.000 val_acc:
0.976 (7.88s - 9.40s remaining)
EPOCH: 229.000 train_loss: 0.010 val_loss: 0.088 train_acc: 1.000 val_acc:
0.976 (7.91s - 9.36s remaining)
EPOCH: 230.000 train_loss: 0.010 val_loss: 0.088 train_acc: 1.000 val_acc:
0.976 (7.94s - 9.32s remaining)
EPOCH: 231.000 train_loss: 0.010 val_loss: 0.088 train_acc: 1.000 val_acc:
0.977 (7.98s - 9.30s remaining)
EPOCH: 232.000 train_loss: 0.010 val_loss: 0.088 train_acc: 1.000 val_acc:
0.977 (8.02s - 9.26s remaining)
EPOCH: 233.000 train_loss: 0.010 val_loss: 0.088 train_acc: 1.000 val_acc:
0.977 (8.05s - 9.23s remaining)
EPOCH: 234.000 train_loss: 0.010 val_loss: 0.087 train_acc: 1.000 val_acc:
0.977 (8.09s - 9.20s remaining)
EPOCH: 235.000 train_loss: 0.010 val_loss: 0.086 train_acc: 1.000 val_acc:
0.978 (8.12s - 9.16s remaining)
EPOCH: 236.000 train_loss: 0.010 val_loss: 0.085 train_acc: 1.000 val_acc:
0.978 (8.16s - 9.12s remaining)
EPOCH: 237.000 train_loss: 0.010 val_loss: 0.084 train_acc: 1.000 val_acc:
0.977 (8.19s - 9.09s remaining)
EPOCH: 238.000 train_loss: 0.009 val_loss: 0.084 train_acc: 1.000 val_acc:
0.977 (8.22s - 9.05s remaining)
EPOCH: 239.000 train_loss: 0.009 val_loss: 0.084 train_acc: 1.000 val_acc:

```

0.977 (8.25s - 9.01s remaining)
EPOCH: 240.000 train_loss: 0.009 val_loss: 0.084 train_acc: 1.000 val_acc:
0.978 (8.29s - 8.98s remaining)
EPOCH: 241.000 train_loss: 0.009 val_loss: 0.084 train_acc: 1.000 val_acc:
0.977 (8.32s - 8.95s remaining)
EPOCH: 242.000 train_loss: 0.009 val_loss: 0.085 train_acc: 1.000 val_acc:
0.977 (8.36s - 8.91s remaining)
EPOCH: 243.000 train_loss: 0.009 val_loss: 0.086 train_acc: 1.000 val_acc:
0.977 (8.39s - 8.87s remaining)
EPOCH: 244.000 train_loss: 0.009 val_loss: 0.087 train_acc: 1.000 val_acc:
0.979 (8.42s - 8.84s remaining)
EPOCH: 245.000 train_loss: 0.009 val_loss: 0.088 train_acc: 1.000 val_acc:
0.979 (8.46s - 8.81s remaining)
EPOCH: 246.000 train_loss: 0.009 val_loss: 0.088 train_acc: 1.000 val_acc:
0.978 (8.50s - 8.77s remaining)
EPOCH: 247.000 train_loss: 0.009 val_loss: 0.087 train_acc: 1.000 val_acc:
0.978 (8.53s - 8.74s remaining)
EPOCH: 248.000 train_loss: 0.009 val_loss: 0.085 train_acc: 1.000 val_acc:
0.979 (8.56s - 8.70s remaining)
EPOCH: 249.000 train_loss: 0.009 val_loss: 0.082 train_acc: 1.000 val_acc:
0.979 (8.60s - 8.67s remaining)
EPOCH: 250.000 train_loss: 0.009 val_loss: 0.080 train_acc: 1.000 val_acc:
0.976 (8.63s - 8.63s remaining)
EPOCH: 251.000 train_loss: 0.009 val_loss: 0.080 train_acc: 1.000 val_acc:
0.976 (8.66s - 8.59s remaining)
EPOCH: 252.000 train_loss: 0.008 val_loss: 0.080 train_acc: 1.000 val_acc:
0.978 (8.72s - 8.58s remaining)
EPOCH: 253.000 train_loss: 0.008 val_loss: 0.081 train_acc: 1.000 val_acc:
0.978 (8.76s - 8.55s remaining)
EPOCH: 254.000 train_loss: 0.008 val_loss: 0.084 train_acc: 1.000 val_acc:
0.978 (8.79s - 8.51s remaining)
EPOCH: 255.000 train_loss: 0.008 val_loss: 0.087 train_acc: 1.000 val_acc:
0.980 (8.82s - 8.48s remaining)
EPOCH: 256.000 train_loss: 0.008 val_loss: 0.086 train_acc: 1.000 val_acc:
0.979 (8.85s - 8.44s remaining)
EPOCH: 257.000 train_loss: 0.008 val_loss: 0.081 train_acc: 1.000 val_acc:
0.979 (8.89s - 8.41s remaining)
EPOCH: 258.000 train_loss: 0.008 val_loss: 0.078 train_acc: 1.000 val_acc:
0.977 (8.93s - 8.38s remaining)
EPOCH: 259.000 train_loss: 0.008 val_loss: 0.077 train_acc: 1.000 val_acc:
0.978 (8.97s - 8.34s remaining)
EPOCH: 260.000 train_loss: 0.008 val_loss: 0.079 train_acc: 1.000 val_acc:
0.978 (9.00s - 8.31s remaining)
EPOCH: 261.000 train_loss: 0.008 val_loss: 0.084 train_acc: 1.000 val_acc:
0.978 (9.03s - 8.27s remaining)
EPOCH: 262.000 train_loss: 0.007 val_loss: 0.083 train_acc: 1.000 val_acc:
0.979 (9.07s - 8.24s remaining)
EPOCH: 263.000 train_loss: 0.007 val_loss: 0.080 train_acc: 1.000 val_acc:

```

```

0.979 (9.12s - 8.22s remaining)
EPOCH: 264.000 train_loss: 0.007 val_loss: 0.077 train_acc: 1.000 val_acc:
0.978 (9.17s - 8.20s remaining)
EPOCH: 265.000 train_loss: 0.007 val_loss: 0.077 train_acc: 1.000 val_acc:
0.978 (9.21s - 8.17s remaining)
EPOCH: 266.000 train_loss: 0.007 val_loss: 0.080 train_acc: 1.000 val_acc:
0.977 (9.26s - 8.15s remaining)
EPOCH: 267.000 train_loss: 0.007 val_loss: 0.082 train_acc: 1.000 val_acc:
0.978 (9.32s - 8.13s remaining)
EPOCH: 268.000 train_loss: 0.007 val_loss: 0.080 train_acc: 1.000 val_acc:
0.979 (9.36s - 8.10s remaining)
EPOCH: 269.000 train_loss: 0.007 val_loss: 0.078 train_acc: 1.000 val_acc:
0.980 (9.40s - 8.08s remaining)
EPOCH: 270.000 train_loss: 0.007 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (9.45s - 8.05s remaining)
EPOCH: 271.000 train_loss: 0.007 val_loss: 0.078 train_acc: 1.000 val_acc:
0.977 (9.49s - 8.02s remaining)
EPOCH: 272.000 train_loss: 0.007 val_loss: 0.080 train_acc: 1.000 val_acc:
0.977 (9.54s - 7.99s remaining)
EPOCH: 273.000 train_loss: 0.007 val_loss: 0.080 train_acc: 1.000 val_acc:
0.979 (9.58s - 7.96s remaining)
EPOCH: 274.000 train_loss: 0.007 val_loss: 0.079 train_acc: 1.000 val_acc:
0.979 (9.62s - 7.94s remaining)
EPOCH: 275.000 train_loss: 0.006 val_loss: 0.078 train_acc: 1.000 val_acc:
0.980 (9.66s - 7.91s remaining)
EPOCH: 276.000 train_loss: 0.006 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (9.71s - 7.88s remaining)
EPOCH: 277.000 train_loss: 0.006 val_loss: 0.078 train_acc: 1.000 val_acc:
0.977 (9.76s - 7.86s remaining)
EPOCH: 278.000 train_loss: 0.006 val_loss: 0.079 train_acc: 1.000 val_acc:
0.977 (9.81s - 7.83s remaining)
EPOCH: 279.000 train_loss: 0.006 val_loss: 0.079 train_acc: 1.000 val_acc:
0.979 (9.86s - 7.81s remaining)
EPOCH: 280.000 train_loss: 0.006 val_loss: 0.079 train_acc: 1.000 val_acc:
0.979 (9.90s - 7.78s remaining)
EPOCH: 281.000 train_loss: 0.006 val_loss: 0.078 train_acc: 1.000 val_acc:
0.980 (9.96s - 7.76s remaining)
EPOCH: 282.000 train_loss: 0.006 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (10.00s - 7.73s remaining)
EPOCH: 283.000 train_loss: 0.006 val_loss: 0.077 train_acc: 1.000 val_acc:
0.978 (10.04s - 7.70s remaining)
EPOCH: 284.000 train_loss: 0.006 val_loss: 0.078 train_acc: 1.000 val_acc:
0.977 (10.08s - 7.67s remaining)
EPOCH: 285.000 train_loss: 0.006 val_loss: 0.078 train_acc: 1.000 val_acc:
0.978 (10.14s - 7.65s remaining)
EPOCH: 286.000 train_loss: 0.006 val_loss: 0.078 train_acc: 1.000 val_acc:
0.979 (10.18s - 7.62s remaining)
EPOCH: 287.000 train_loss: 0.006 val_loss: 0.078 train_acc: 1.000 val_acc:

```

```

0.979 (10.23s - 7.59s remaining)
EPOCH: 288.000 train_loss: 0.006 val_loss: 0.077 train_acc: 1.000 val_acc:
0.980 (10.27s - 7.56s remaining)
EPOCH: 289.000 train_loss: 0.006 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (10.31s - 7.53s remaining)
EPOCH: 290.000 train_loss: 0.006 val_loss: 0.077 train_acc: 1.000 val_acc:
0.978 (10.37s - 7.51s remaining)
EPOCH: 291.000 train_loss: 0.006 val_loss: 0.077 train_acc: 1.000 val_acc:
0.977 (10.41s - 7.48s remaining)
EPOCH: 292.000 train_loss: 0.006 val_loss: 0.078 train_acc: 1.000 val_acc:
0.979 (10.46s - 7.45s remaining)
EPOCH: 293.000 train_loss: 0.005 val_loss: 0.078 train_acc: 1.000 val_acc:
0.979 (10.51s - 7.42s remaining)
EPOCH: 294.000 train_loss: 0.005 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (10.56s - 7.40s remaining)
EPOCH: 295.000 train_loss: 0.005 val_loss: 0.077 train_acc: 1.000 val_acc:
0.980 (10.60s - 7.37s remaining)
EPOCH: 296.000 train_loss: 0.005 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (10.64s - 7.33s remaining)
EPOCH: 297.000 train_loss: 0.005 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (10.69s - 7.30s remaining)
EPOCH: 298.000 train_loss: 0.005 val_loss: 0.076 train_acc: 1.000 val_acc:
0.978 (10.73s - 7.27s remaining)
EPOCH: 299.000 train_loss: 0.005 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (10.78s - 7.25s remaining)
EPOCH: 300.000 train_loss: 0.005 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (10.84s - 7.22s remaining)
EPOCH: 301.000 train_loss: 0.005 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (10.88s - 7.20s remaining)
EPOCH: 302.000 train_loss: 0.005 val_loss: 0.077 train_acc: 1.000 val_acc:
0.980 (10.93s - 7.16s remaining)
EPOCH: 303.000 train_loss: 0.005 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (10.98s - 7.14s remaining)
EPOCH: 304.000 train_loss: 0.005 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (11.02s - 7.11s remaining)
EPOCH: 305.000 train_loss: 0.005 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (11.07s - 7.08s remaining)
EPOCH: 306.000 train_loss: 0.005 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (11.12s - 7.05s remaining)
EPOCH: 307.000 train_loss: 0.005 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (11.18s - 7.03s remaining)
EPOCH: 308.000 train_loss: 0.005 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (11.23s - 7.00s remaining)
EPOCH: 309.000 train_loss: 0.005 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (11.28s - 6.97s remaining)
EPOCH: 310.000 train_loss: 0.005 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (11.33s - 6.94s remaining)
EPOCH: 311.000 train_loss: 0.005 val_loss: 0.076 train_acc: 1.000 val_acc:

```



```

0.980 (11.38s - 6.92s remaining)
EPOCH: 312.000 train_loss: 0.005 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (11.43s - 6.89s remaining)
EPOCH: 313.000 train_loss: 0.005 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (11.47s - 6.85s remaining)
EPOCH: 314.000 train_loss: 0.005 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (11.52s - 6.82s remaining)
EPOCH: 315.000 train_loss: 0.005 val_loss: 0.075 train_acc: 1.000 val_acc:
0.979 (11.56s - 6.79s remaining)
EPOCH: 316.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (11.61s - 6.76s remaining)
EPOCH: 317.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (11.66s - 6.73s remaining)
EPOCH: 318.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (11.70s - 6.70s remaining)
EPOCH: 319.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (11.75s - 6.67s remaining)
EPOCH: 320.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (11.80s - 6.64s remaining)
EPOCH: 321.000 train_loss: 0.004 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (11.85s - 6.61s remaining)
EPOCH: 322.000 train_loss: 0.004 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (11.90s - 6.58s remaining)
EPOCH: 323.000 train_loss: 0.004 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (11.95s - 6.55s remaining)
EPOCH: 324.000 train_loss: 0.004 val_loss: 0.074 train_acc: 1.000 val_acc:
0.979 (11.99s - 6.52s remaining)
EPOCH: 325.000 train_loss: 0.004 val_loss: 0.075 train_acc: 1.000 val_acc:
0.979 (12.03s - 6.48s remaining)
EPOCH: 326.000 train_loss: 0.004 val_loss: 0.075 train_acc: 1.000 val_acc:
0.979 (12.06s - 6.44s remaining)
EPOCH: 327.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (12.10s - 6.40s remaining)
EPOCH: 328.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (12.13s - 6.36s remaining)
EPOCH: 329.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (12.17s - 6.33s remaining)
EPOCH: 330.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (12.21s - 6.29s remaining)
EPOCH: 331.000 train_loss: 0.004 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (12.24s - 6.25s remaining)
EPOCH: 332.000 train_loss: 0.004 val_loss: 0.073 train_acc: 1.000 val_acc:
0.978 (12.28s - 6.21s remaining)
EPOCH: 333.000 train_loss: 0.004 val_loss: 0.073 train_acc: 1.000 val_acc:
0.978 (12.31s - 6.17s remaining)
EPOCH: 334.000 train_loss: 0.004 val_loss: 0.073 train_acc: 1.000 val_acc:
0.977 (12.34s - 6.13s remaining)
EPOCH: 335.000 train_loss: 0.004 val_loss: 0.074 train_acc: 1.000 val_acc:

```

```

0.979 (12.38s - 6.10s remaining)
EPOCH: 336.000 train_loss: 0.004 val_loss: 0.075 train_acc: 1.000 val_acc:
0.978 (12.42s - 6.06s remaining)
EPOCH: 337.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (12.45s - 6.02s remaining)
EPOCH: 338.000 train_loss: 0.004 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (12.48s - 5.98s remaining)
EPOCH: 339.000 train_loss: 0.004 val_loss: 0.077 train_acc: 1.000 val_acc:
0.980 (12.51s - 5.94s remaining)
EPOCH: 340.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.981 (12.55s - 5.90s remaining)
EPOCH: 341.000 train_loss: 0.004 val_loss: 0.075 train_acc: 1.000 val_acc:
0.981 (12.59s - 5.87s remaining)
EPOCH: 342.000 train_loss: 0.004 val_loss: 0.073 train_acc: 1.000 val_acc:
0.978 (12.63s - 5.83s remaining)
EPOCH: 343.000 train_loss: 0.004 val_loss: 0.072 train_acc: 1.000 val_acc:
0.977 (12.66s - 5.80s remaining)
EPOCH: 344.000 train_loss: 0.004 val_loss: 0.071 train_acc: 1.000 val_acc:
0.976 (12.69s - 5.76s remaining)
EPOCH: 345.000 train_loss: 0.004 val_loss: 0.072 train_acc: 1.000 val_acc:
0.978 (12.73s - 5.72s remaining)
EPOCH: 346.000 train_loss: 0.004 val_loss: 0.074 train_acc: 1.000 val_acc:
0.978 (12.76s - 5.68s remaining)
EPOCH: 347.000 train_loss: 0.004 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (12.79s - 5.64s remaining)
EPOCH: 348.000 train_loss: 0.004 val_loss: 0.078 train_acc: 1.000 val_acc:
0.981 (12.84s - 5.61s remaining)
EPOCH: 349.000 train_loss: 0.004 val_loss: 0.078 train_acc: 1.000 val_acc:
0.980 (12.87s - 5.57s remaining)
EPOCH: 350.000 train_loss: 0.004 val_loss: 0.076 train_acc: 1.000 val_acc:
0.981 (12.91s - 5.53s remaining)
EPOCH: 351.000 train_loss: 0.004 val_loss: 0.073 train_acc: 1.000 val_acc:
0.979 (12.94s - 5.49s remaining)
EPOCH: 352.000 train_loss: 0.003 val_loss: 0.071 train_acc: 1.000 val_acc:
0.978 (12.98s - 5.46s remaining)
EPOCH: 353.000 train_loss: 0.003 val_loss: 0.069 train_acc: 1.000 val_acc:
0.979 (13.02s - 5.42s remaining)
EPOCH: 354.000 train_loss: 0.003 val_loss: 0.071 train_acc: 1.000 val_acc:
0.979 (13.05s - 5.38s remaining)
EPOCH: 355.000 train_loss: 0.003 val_loss: 0.074 train_acc: 1.000 val_acc:
0.977 (13.09s - 5.34s remaining)
EPOCH: 356.000 train_loss: 0.003 val_loss: 0.078 train_acc: 1.000 val_acc:
0.980 (13.13s - 5.31s remaining)
EPOCH: 357.000 train_loss: 0.003 val_loss: 0.078 train_acc: 1.000 val_acc:
0.980 (13.16s - 5.27s remaining)
EPOCH: 358.000 train_loss: 0.003 val_loss: 0.075 train_acc: 1.000 val_acc:
0.981 (13.20s - 5.24s remaining)
EPOCH: 359.000 train_loss: 0.003 val_loss: 0.072 train_acc: 1.000 val_acc:

```

0.978 (13.24s - 5.20s remaining)
EPOCH: 360.000 train_loss: 0.003 val_loss: 0.069 train_acc: 1.000 val_acc:
0.980 (13.28s - 5.16s remaining)
EPOCH: 361.000 train_loss: 0.003 val_loss: 0.070 train_acc: 1.000 val_acc:
0.979 (13.31s - 5.13s remaining)
EPOCH: 362.000 train_loss: 0.003 val_loss: 0.075 train_acc: 1.000 val_acc:
0.977 (13.35s - 5.09s remaining)
EPOCH: 363.000 train_loss: 0.003 val_loss: 0.077 train_acc: 1.000 val_acc:
0.981 (13.38s - 5.05s remaining)
EPOCH: 364.000 train_loss: 0.003 val_loss: 0.075 train_acc: 1.000 val_acc:
0.981 (13.41s - 5.01s remaining)
EPOCH: 365.000 train_loss: 0.003 val_loss: 0.072 train_acc: 1.000 val_acc:
0.979 (13.45s - 4.98s remaining)
EPOCH: 366.000 train_loss: 0.003 val_loss: 0.070 train_acc: 1.000 val_acc:
0.979 (13.49s - 4.94s remaining)
EPOCH: 367.000 train_loss: 0.003 val_loss: 0.071 train_acc: 1.000 val_acc:
0.981 (13.54s - 4.91s remaining)
EPOCH: 368.000 train_loss: 0.003 val_loss: 0.074 train_acc: 1.000 val_acc:
0.979 (13.59s - 4.87s remaining)
EPOCH: 369.000 train_loss: 0.003 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (13.64s - 4.84s remaining)
EPOCH: 370.000 train_loss: 0.003 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (13.69s - 4.81s remaining)
EPOCH: 371.000 train_loss: 0.003 val_loss: 0.073 train_acc: 1.000 val_acc:
0.979 (13.73s - 4.78s remaining)
EPOCH: 372.000 train_loss: 0.003 val_loss: 0.071 train_acc: 1.000 val_acc:
0.978 (13.79s - 4.74s remaining)
EPOCH: 373.000 train_loss: 0.003 val_loss: 0.071 train_acc: 1.000 val_acc:
0.980 (13.86s - 4.72s remaining)
EPOCH: 374.000 train_loss: 0.003 val_loss: 0.073 train_acc: 1.000 val_acc:
0.979 (13.92s - 4.69s remaining)
EPOCH: 375.000 train_loss: 0.003 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (13.97s - 4.66s remaining)
EPOCH: 376.000 train_loss: 0.003 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (14.00s - 4.62s remaining)
EPOCH: 377.000 train_loss: 0.003 val_loss: 0.073 train_acc: 1.000 val_acc:
0.980 (14.03s - 4.58s remaining)
EPOCH: 378.000 train_loss: 0.003 val_loss: 0.072 train_acc: 1.000 val_acc:
0.978 (14.07s - 4.54s remaining)
EPOCH: 379.000 train_loss: 0.003 val_loss: 0.072 train_acc: 1.000 val_acc:
0.979 (14.11s - 4.50s remaining)
EPOCH: 380.000 train_loss: 0.003 val_loss: 0.073 train_acc: 1.000 val_acc:
0.979 (14.14s - 4.47s remaining)
EPOCH: 381.000 train_loss: 0.003 val_loss: 0.074 train_acc: 1.000 val_acc:
0.979 (14.18s - 4.43s remaining)
EPOCH: 382.000 train_loss: 0.003 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (14.21s - 4.39s remaining)
EPOCH: 383.000 train_loss: 0.003 val_loss: 0.074 train_acc: 1.000 val_acc:

```

0.981 (14.25s - 4.35s remaining)
EPOCH: 384.000 train_loss: 0.003 val_loss: 0.073 train_acc: 1.000 val_acc:
0.980 (14.29s - 4.32s remaining)
EPOCH: 385.000 train_loss: 0.003 val_loss: 0.072 train_acc: 1.000 val_acc:
0.978 (14.32s - 4.28s remaining)
EPOCH: 386.000 train_loss: 0.003 val_loss: 0.072 train_acc: 1.000 val_acc:
0.979 (14.36s - 4.24s remaining)
EPOCH: 387.000 train_loss: 0.003 val_loss: 0.073 train_acc: 1.000 val_acc:
0.979 (14.40s - 4.20s remaining)
EPOCH: 388.000 train_loss: 0.003 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (14.43s - 4.17s remaining)
EPOCH: 389.000 train_loss: 0.003 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (14.47s - 4.13s remaining)
EPOCH: 390.000 train_loss: 0.003 val_loss: 0.074 train_acc: 1.000 val_acc:
0.981 (14.51s - 4.09s remaining)
EPOCH: 391.000 train_loss: 0.003 val_loss: 0.073 train_acc: 1.000 val_acc:
0.980 (14.54s - 4.05s remaining)
EPOCH: 392.000 train_loss: 0.003 val_loss: 0.072 train_acc: 1.000 val_acc:
0.979 (14.57s - 4.02s remaining)
EPOCH: 393.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.979 (14.61s - 3.98s remaining)
EPOCH: 394.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.980 (14.64s - 3.94s remaining)
EPOCH: 395.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.979 (14.67s - 3.90s remaining)
EPOCH: 396.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (14.71s - 3.86s remaining)
EPOCH: 397.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (14.74s - 3.83s remaining)
EPOCH: 398.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.981 (14.78s - 3.79s remaining)
EPOCH: 399.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.980 (14.82s - 3.75s remaining)
EPOCH: 400.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.979 (14.85s - 3.71s remaining)
EPOCH: 401.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.979 (14.89s - 3.67s remaining)
EPOCH: 402.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.980 (14.92s - 3.64s remaining)
EPOCH: 403.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (14.96s - 3.60s remaining)
EPOCH: 404.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (15.00s - 3.56s remaining)
EPOCH: 405.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (15.03s - 3.53s remaining)
EPOCH: 406.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.981 (15.07s - 3.49s remaining)
EPOCH: 407.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:

```

```

0.980 (15.11s - 3.45s remaining)
EPOCH: 408.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.979 (15.14s - 3.41s remaining)
EPOCH: 409.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.979 (15.18s - 3.38s remaining)
EPOCH: 410.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.980 (15.21s - 3.34s remaining)
EPOCH: 411.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.979 (15.24s - 3.30s remaining)
EPOCH: 412.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (15.28s - 3.26s remaining)
EPOCH: 413.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (15.33s - 3.23s remaining)
EPOCH: 414.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (15.36s - 3.19s remaining)
EPOCH: 415.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (15.39s - 3.15s remaining)
EPOCH: 416.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.979 (15.43s - 3.11s remaining)
EPOCH: 417.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.978 (15.46s - 3.08s remaining)
EPOCH: 418.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.979 (15.49s - 3.04s remaining)
EPOCH: 419.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.979 (15.54s - 3.00s remaining)
EPOCH: 420.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (15.58s - 2.97s remaining)
EPOCH: 421.000 train_loss: 0.002 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (15.61s - 2.93s remaining)
EPOCH: 422.000 train_loss: 0.002 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (15.65s - 2.89s remaining)
EPOCH: 423.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.981 (15.68s - 2.85s remaining)
EPOCH: 424.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (15.71s - 2.82s remaining)
EPOCH: 425.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.980 (15.75s - 2.78s remaining)
EPOCH: 426.000 train_loss: 0.002 val_loss: 0.071 train_acc: 1.000 val_acc:
0.980 (15.79s - 2.74s remaining)
EPOCH: 427.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.982 (15.82s - 2.70s remaining)
EPOCH: 428.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.980 (15.86s - 2.67s remaining)
EPOCH: 429.000 train_loss: 0.002 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (15.90s - 2.63s remaining)
EPOCH: 430.000 train_loss: 0.002 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (15.94s - 2.59s remaining)
EPOCH: 431.000 train_loss: 0.002 val_loss: 0.077 train_acc: 1.000 val_acc:

```

```

0.979 (15.98s - 2.56s remaining)
EPOCH: 432.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.981 (16.01s - 2.52s remaining)
EPOCH: 433.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.979 (16.04s - 2.48s remaining)
EPOCH: 434.000 train_loss: 0.002 val_loss: 0.071 train_acc: 1.000 val_acc:
0.980 (16.08s - 2.44s remaining)
EPOCH: 435.000 train_loss: 0.002 val_loss: 0.070 train_acc: 1.000 val_acc:
0.980 (16.11s - 2.41s remaining)
EPOCH: 436.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.981 (16.15s - 2.37s remaining)
EPOCH: 437.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.979 (16.19s - 2.33s remaining)
EPOCH: 438.000 train_loss: 0.002 val_loss: 0.078 train_acc: 1.000 val_acc:
0.979 (16.22s - 2.30s remaining)
EPOCH: 439.000 train_loss: 0.002 val_loss: 0.078 train_acc: 1.000 val_acc:
0.979 (16.25s - 2.26s remaining)
EPOCH: 440.000 train_loss: 0.002 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (16.29s - 2.22s remaining)
EPOCH: 441.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.980 (16.32s - 2.18s remaining)
EPOCH: 442.000 train_loss: 0.002 val_loss: 0.071 train_acc: 1.000 val_acc:
0.980 (16.37s - 2.15s remaining)
EPOCH: 443.000 train_loss: 0.002 val_loss: 0.070 train_acc: 1.000 val_acc:
0.981 (16.40s - 2.11s remaining)
EPOCH: 444.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.981 (16.44s - 2.07s remaining)
EPOCH: 445.000 train_loss: 0.002 val_loss: 0.077 train_acc: 1.000 val_acc:
0.980 (16.47s - 2.04s remaining)
EPOCH: 446.000 train_loss: 0.002 val_loss: 0.078 train_acc: 1.000 val_acc:
0.979 (16.51s - 2.00s remaining)
EPOCH: 447.000 train_loss: 0.002 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (16.55s - 1.96s remaining)
EPOCH: 448.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.979 (16.58s - 1.92s remaining)
EPOCH: 449.000 train_loss: 0.002 val_loss: 0.070 train_acc: 1.000 val_acc:
0.980 (16.61s - 1.89s remaining)
EPOCH: 450.000 train_loss: 0.002 val_loss: 0.070 train_acc: 1.000 val_acc:
0.981 (16.65s - 1.85s remaining)
EPOCH: 451.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.981 (16.68s - 1.81s remaining)
EPOCH: 452.000 train_loss: 0.002 val_loss: 0.077 train_acc: 1.000 val_acc:
0.979 (16.71s - 1.77s remaining)
EPOCH: 453.000 train_loss: 0.002 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (16.75s - 1.74s remaining)
EPOCH: 454.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (16.79s - 1.70s remaining)
EPOCH: 455.000 train_loss: 0.002 val_loss: 0.071 train_acc: 1.000 val_acc:

```

```

0.980 (16.82s - 1.66s remaining)
EPOCH: 456.000 train_loss: 0.002 val_loss: 0.071 train_acc: 1.000 val_acc:
0.981 (16.86s - 1.63s remaining)
EPOCH: 457.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.982 (16.90s - 1.59s remaining)
EPOCH: 458.000 train_loss: 0.002 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (16.93s - 1.55s remaining)
EPOCH: 459.000 train_loss: 0.002 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (16.98s - 1.52s remaining)
EPOCH: 460.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.981 (17.02s - 1.48s remaining)
EPOCH: 461.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.982 (17.05s - 1.44s remaining)
EPOCH: 462.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.983 (17.09s - 1.41s remaining)
EPOCH: 463.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.982 (17.12s - 1.37s remaining)
EPOCH: 464.000 train_loss: 0.002 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (17.16s - 1.33s remaining)
EPOCH: 465.000 train_loss: 0.002 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (17.20s - 1.29s remaining)
EPOCH: 466.000 train_loss: 0.002 val_loss: 0.074 train_acc: 1.000 val_acc:
0.982 (17.23s - 1.26s remaining)
EPOCH: 467.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.981 (17.27s - 1.22s remaining)
EPOCH: 468.000 train_loss: 0.002 val_loss: 0.072 train_acc: 1.000 val_acc:
0.983 (17.30s - 1.18s remaining)
EPOCH: 469.000 train_loss: 0.002 val_loss: 0.073 train_acc: 1.000 val_acc:
0.982 (17.34s - 1.15s remaining)
EPOCH: 470.000 train_loss: 0.001 val_loss: 0.074 train_acc: 1.000 val_acc:
0.980 (17.38s - 1.11s remaining)
EPOCH: 471.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (17.42s - 1.07s remaining)
EPOCH: 472.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (17.45s - 1.04s remaining)
EPOCH: 473.000 train_loss: 0.001 val_loss: 0.074 train_acc: 1.000 val_acc:
0.982 (17.48s - 1.00s remaining)
EPOCH: 474.000 train_loss: 0.001 val_loss: 0.073 train_acc: 1.000 val_acc:
0.984 (17.52s - 0.96s remaining)
EPOCH: 475.000 train_loss: 0.001 val_loss: 0.072 train_acc: 1.000 val_acc:
0.984 (17.55s - 0.92s remaining)
EPOCH: 476.000 train_loss: 0.001 val_loss: 0.074 train_acc: 1.000 val_acc:
0.982 (17.59s - 0.89s remaining)
EPOCH: 477.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (17.63s - 0.85s remaining)
EPOCH: 478.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (17.66s - 0.81s remaining)
EPOCH: 479.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:

```

```

0.981 (17.69s - 0.78s remaining)
EPOCH: 480.000 train_loss: 0.001 val_loss: 0.073 train_acc: 1.000 val_acc:
0.982 (17.72s - 0.74s remaining)
EPOCH: 481.000 train_loss: 0.001 val_loss: 0.073 train_acc: 1.000 val_acc:
0.984 (17.77s - 0.70s remaining)
EPOCH: 482.000 train_loss: 0.001 val_loss: 0.073 train_acc: 1.000 val_acc:
0.984 (17.81s - 0.67s remaining)
EPOCH: 483.000 train_loss: 0.001 val_loss: 0.074 train_acc: 1.000 val_acc:
0.982 (17.84s - 0.63s remaining)
EPOCH: 484.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (17.88s - 0.59s remaining)
EPOCH: 485.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (17.91s - 0.55s remaining)
EPOCH: 486.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (17.94s - 0.52s remaining)
EPOCH: 487.000 train_loss: 0.001 val_loss: 0.074 train_acc: 1.000 val_acc:
0.982 (17.98s - 0.48s remaining)
EPOCH: 488.000 train_loss: 0.001 val_loss: 0.073 train_acc: 1.000 val_acc:
0.984 (18.03s - 0.44s remaining)
EPOCH: 489.000 train_loss: 0.001 val_loss: 0.073 train_acc: 1.000 val_acc:
0.984 (18.06s - 0.41s remaining)
EPOCH: 490.000 train_loss: 0.001 val_loss: 0.074 train_acc: 1.000 val_acc:
0.983 (18.10s - 0.37s remaining)
EPOCH: 491.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (18.14s - 0.33s remaining)
EPOCH: 492.000 train_loss: 0.001 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (18.17s - 0.30s remaining)
EPOCH: 493.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:
0.980 (18.22s - 0.26s remaining)
EPOCH: 494.000 train_loss: 0.001 val_loss: 0.074 train_acc: 1.000 val_acc:
0.981 (18.26s - 0.22s remaining)
EPOCH: 495.000 train_loss: 0.001 val_loss: 0.073 train_acc: 1.000 val_acc:
0.984 (18.29s - 0.18s remaining)
EPOCH: 496.000 train_loss: 0.001 val_loss: 0.073 train_acc: 1.000 val_acc:
0.984 (18.32s - 0.15s remaining)
EPOCH: 497.000 train_loss: 0.001 val_loss: 0.073 train_acc: 1.000 val_acc:
0.983 (18.36s - 0.11s remaining)
EPOCH: 498.000 train_loss: 0.001 val_loss: 0.075 train_acc: 1.000 val_acc:
0.981 (18.39s - 0.07s remaining)
EPOCH: 499.000 train_loss: 0.001 val_loss: 0.076 train_acc: 1.000 val_acc:
0.980 (18.43s - 0.04s remaining)
EPOCH: 500.000 train_loss: 0.001 val_loss: 0.076 train_acc: 1.000 val_acc:
0.979 (18.47s - 0.00s remaining)

```

```

[48]: lstm_results= {"train_acc": train_accuracy, "val_acc": val_accuracy,
↳ "train_loss": train_loss, "val_loss": val_loss, "epochs": 5}

```


##2.4. Plotting the loss and acc curves


```
[49]: fig, ax= plt.subplots(ncols=2, figsize=(15,5)) #w,h
log.plot_epochs(["train_loss", "val_loss"], ax=ax[0], title="Loss_curve");
    ↪ #loss_curve
ax[0].legend(loc='upper right', fontsize=10)
log.plot_epochs(["train_acc", "val_acc"], ax=ax[1], title="Acc_curve");
    ↪ #acc_curve
```

100%| | 601/601 [00:00<00:00, 1450.76it/s]

WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

100%| | 601/601 [00:00<00:00, 1257.36it/s]



##2.5. Testing on a new datapoint

```
[50]: x=gen_bin(10)
inp, tar= input_prep(x[45], x[64])
x[45], x[64], inp, tar
```

```
[50]: ('0b10000101101',
       '0b10001000000',
       tensor([[1., 0.],
               [0., 0.],
               [1., 0.],
               [1., 0.],
               [0., 0.],
               [1., 0.],
               [0., 1.],
               [0., 0.],
               [0., 0.],
               [0., 0.]]),
       tensor([1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0]))
```

```
[51]: inp=torch.unsqueeze(inp, dim=0)
inp.shape
```

```
[51]: torch.Size([1, 10, 2])
```

```
[54]: torch.round(lstm(inp.to(device)))
```

```
[54]: tensor([[1., 0., 1., 1., 1., 0., 1., 0., 0., 0., 0.]], device='cuda:0',
        grad_fn=<RoundBackward0>)
```

We can see that the rnn network performs very well on the test data.

###2.6. Plotting the acc as a function of N

```
[55]: def experiments(lst_len_of_numbers, epochs):
    train_accs, test_accs=[], []
    input_size=2
    hidden_size=64
    num_layers=2
    for N in lst_len_of_numbers:
        train_ds= binary_string_dataset(32*10, N)
        test_ds= binary_string_dataset(32*3, N)
        train_dl= torch.utils.data.DataLoader(train_ds, batch_size=32)
        test_dl= torch.utils.data.DataLoader(test_ds, batch_size=32)

        sequence_length= N
        num_classes= N

        lstm= LSTM(input_size, hidden_size, num_layers, num_classes).to(device)
        EPOCHS=epochs
        criterion= lstm.loss_and_accuracy
        optimizer=torch.optim.Adam(params= lstm.parameters())
        log= Report(EPOCHS)
        for epoch in range(EPOCHS):
            n=len(train_dl)
            for ix, input in enumerate(train_dl):
                train_loss, train_accuracy=train_epoch(lstm, input, criterion,
                optimizer, sequence_length=sequence_length)
                log.record(epoch+(ix+1)/n, train_loss=train_loss, train_acc=
                train_accuracy, end="\r")
            n=len(test_dl)
            for ix, input in enumerate(test_dl):
                val_loss, val_accuracy=val_epoch(lstm, input, criterion,
                sequence_length=sequence_length)
                log.record(epoch+(ix+1)/n, val_loss=val_loss, val_acc= val_accuracy,
                end="\r")
            train_accs.append(train_accuracy)
            test_accs.append(val_accuracy)
```

```
return train_accs, test_accs
```

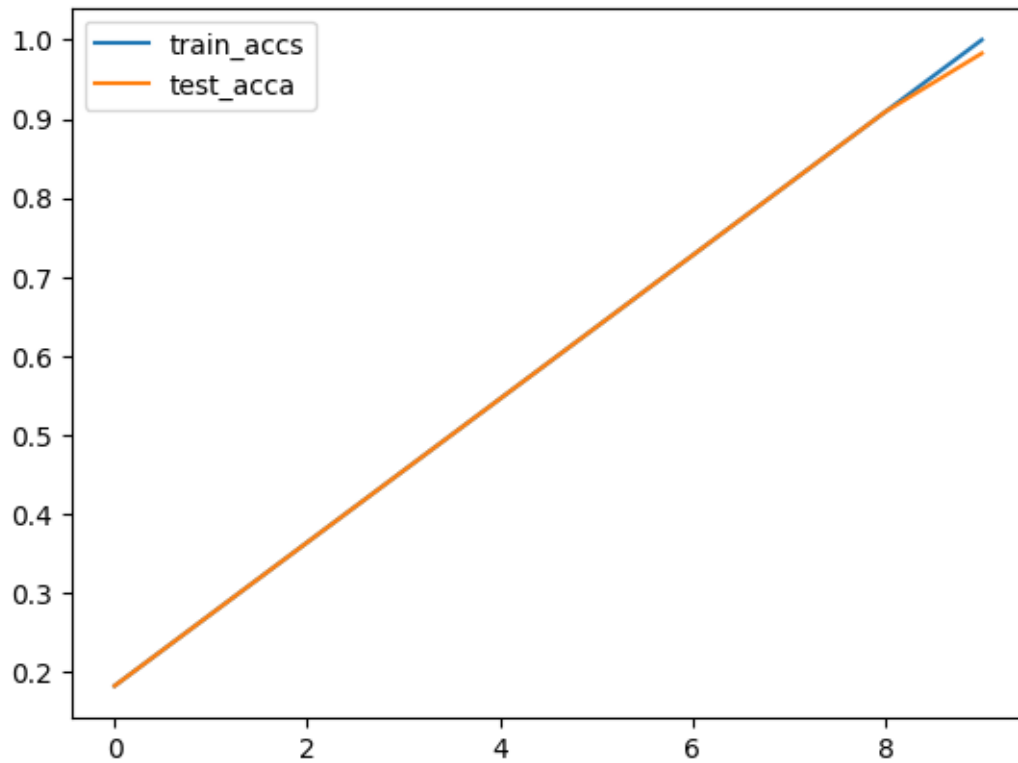
```
[56]: lst_len_of_numbers=np.arange(1, 11)
epochs=500
train_accs, test_accs= experiments(lst_len_of_numbers, epochs)
```

EPOCH: 500.000 val_loss: 0.098 val_acc: 0.983 (17.48s - 0.00s remaining)

```
[57]: train_accs, test_accs
```

```
[57]: ([0.1818181872367859,
0.27272728085517883,
0.3636363744735718,
0.4545454680919647,
0.5454545617103577,
0.6363636255264282,
0.7272727489471436,
0.8181818723678589,
0.9090909361839294,
1.0],
[0.1818181872367859,
0.27272728085517883,
0.3636363744735718,
0.4545454680919647,
0.5454545617103577,
0.6363636255264282,
0.7272727489471436,
0.8181818723678589,
0.9090909361839294,
0.9829545617103577])
```

```
[58]: plt.plot(train_accs, label="train_accs")
plt.plot(test_accs, label="test_acca")
plt.legend()
plt.show();
```



For the given hyperparameters we see a linearly increasing graph, it can also be seen that the test_accuracy starts lagging near the tip end.